Round 001 reviewer-response addendum

R1.1 — Practical appearance. The main pilot table now includes the measured GT-256 reference bake, with the same source resolution, masks and baking settings. Its mean FG-LPIPS is 0.0590, compared with 0.1303 for Fixed Mean and 0.1277 for Late-high. The named Late-high gain is stable in direction but small in context: a post hoc ratio of paired means is 3.7% of the Fixed Mean–GT reference gap (pointwise object-bootstrap CI 2.3–5.3%). The empirical reference is not a theoretical bound or a material-quality threshold. The substantial appearance failure remains visible. We do not infer acceptable PBR fidelity from the separated LPIPS contrast.

R1.4 — Uncertainty. The previous frozen primary and its interval are unchanged. The descriptive context resamples paired objects jointly; it is not added to the confirmatory test family. Leaving out each object keeps the mean difference negative, between -0.00294 and -0.00208, but does not supply new independent evidence or an equivalence test.

R2.1 — Contribution. This round improves the interpretation and evidence completeness; it does not add a new scheduling algorithm or claim to resolve novelty. No acceptance probability is inferred.