Round 002 reviewer-response addendum

R1.1 — Default controls and practical appearance. We added the actual unset-scale GeoTex wrapper residual and No Adapter on the same ten previously inspected pilot objects. The regenerated capped 1.0 control reproduces all ten previous RGB outputs bitwise; new controls preserve mesh bytes, camera/mask settings and raw atlas coverage. Repeated GT endpoints match exactly. The native residual branch has unseen-view FG-PSNR 11.13 dB, FG-LPIPS 0.1372, FG-SSIM 0.3376 and CIEDE2000 26.47. No Adapter has 5.92 dB, 0.1600, 0.3829 and 40.92, respectively. All valid failures remain displayed in complete six-view grids.

These are controls inside the frozen 50-step sampling protocol, not a replay of the complete default pipeline. A read-only source audit establishes that the frozen loop uses 256-pixel views, cached image embeddings and conditional-only score predictions; the full pipeline defaults to 512-pixel views, online image embeddings and classifier-free guidance 2.0. The model validation call uses 75 steps and does not explicitly bind geometry features. We state this distinction in the methods and Supplement S20. It does not by itself identify the cause of poor appearance or validate a repair.

R1.4 — Complete statistics. The added analyses are explicitly post hoc, with three named comparisons on five endpoints in one 15-test Holm family. Late-high–native residual PSNR is +1.79 dB (pointwise 95% object-bootstrap CI +0.46 to +3.06), and LPIPS is −0.00951 (−0.01772 to −0.00263); the respective adjusted p values are both 0.297. C3–native residual has +1.35 dB and−0.00699, also adjusted p 0.297. None of the ten schedule contrasts separates after correction; only No Adapter–native residual Edge-SSIM does (−0.218; adjusted p 0.0293). The older primary, cohort, strong-baseline results and GT reference context remain unchanged. Non-rejection is not equivalence.

R2.1 — Contribution boundary and next decision. This round closes the capped/unset branch ambiguity within the tested sampling stack; it does not establish full-pipeline equivalence, a new algorithm or generally acceptable baked appearance. Complete default-path quality now has higher priority than another schedule search. Novelty remains a substantive limitation, and no acceptance probability is inferred.