## Seventh-round response addendum

We separate geometry-encoding size (256/512) and CFG correction policy (both, conditional-only, conditional-only half amplitude) in twelve seen-case native-loop calls. The256/both arm reproduces both preceding original-residual grids bitwise; every conditional-only correction records a zero negative branch. Inputs/noise match across arms and geometry features match within encoding size. Half-amplitude conditional correction retains more tire detail, but neither factor restores the remote reference content. All outputs and negative findings are retained in S25.

The recorded custom training entry uses conditional inputs throughout and does not read the conditioning-drop probability declared in its model configuration. This entry differs from the model-class training method. The audit also records256 training, historical repeated-top targets, checkpoint/list/source hashes and complete2,000-step loss rows. The1,118-object training list has no Fresh C UID overlap;1,112 have the six checked input sets. Missing seed/list-override records and the changed surviving dataset source prevent an exact-replay claim. These observations are not restoration of historical FAC or causal proof of the failure. Removing the unconditioned correction alone does not repair appearance. Protocol-consistent training checks now take priority over further schedule sweeps; independent native deployment benefits and novelty remain unresolved.