Reviewer risks and experiment decisions

10 October 2026 · Author review only · `SUBMISSION_APPROVED=NO` · No new GPU run launched

Author checks: compare the response with the exact second-round decision; confirm Figure 6 rights or replace it.

Issue	Evidence and remaining risk	Decision
R1.1 Color / 3D fidelity — high, partial	Failure-enriched N=26 diagnosis; no validated color repair. No Adapter is not native-default. Fresh C bake gate failed (9/24 out-of-range UV inputs; 5 assets / 27 primitives use unsupported BLEND). Historical N=20 unlit check is mixed. Figure 6 source-asset rights and panel mapping are unverified.	**E1: highest priority, no-go now.** Do a CPU-only preflight on an unchanged, frozen in-cohort asset (estimated 1–2 h). Require valid original UVs/materials, sampler/import-export integrity, finite output, and repeatable pixels. If it passes, a small native-default / No Adapter / GFL / C3 bake and unseen-view/seam pilot is estimated at 0.5–1 day. Failure means stop, no substitution or input repair. Confirm Figure 6 rights or replace before release.
R1.2 / R1.4 Matched result and uncertainty — bounded	Fresh C: N=300; C3−GFL FG-PSNR +0.522352 dB (95% CI [+0.404160,+0.640669]). GFH and Generic Linear have stronger means on all four shown endpoints. Applied deep/middle sums differ, 82.5 vs. 62.5.	No new experiment. Keep the positive endpoint result and counter-results together; no timing-only, equivalence, or non-inferiority claim.
R1.3 Texture fidelity — partial	Variation proxies and subjective preference are not fidelity. Fresh C has no validated reference-texture-fidelity endpoint.	No experiment selected. A future study needs a defined reference, held-out protocol, and validated endpoint before data collection.
R1.5 FAC — partial	Reported training pool is 1,706 objects; recovered config points to 1,200. Seed, output, and runtime are not bound in one manifest.	Retain as prior negative evidence with the mismatch disclosed; make no full-reproduction or public-code claim.
R2.1 Contribution — editorial risk	SSI establishes prior layer/time scheduling, including geometry-conditioned ControlNet strength. TCAS is a task-specific GeoTex-Adapter residual intervention with measured cap, dose, and endpoint limits.	No schedule sweep. Present an empirical graphics-system contribution; contribution sufficiency remains an editor judgment.
R2.3 Cross-backbone — high, open	MV-Adapter (98-object and 76-object probes) and MVDiffusion (75 objects) use different interfaces/runners; corrected MV-Adapter interaction tests have no Holm rejection, which is not equivalence.	**E2: no-go.** First verify that one second backbone can implement the same scale semantics and within-backbone native/fixed/scheduled controls. Desk feasibility: 2–4 h; only if it passes, freeze a small matched pilot (estimated 0.5–1 GPU day). If semantics cannot match, keep interface-specific results and stop.
Dose / temporal placement — open	Strict-276 temporal L–L–H beats C3 on six endpoints, but realized residual norms differ. A3b has no layer-group dose common support (0/2,250 observations).	**E3: no-go.** Reconsider only with a frozen design showing realized-dose overlap and a useful estimand. Protocol/data audit estimate: 0.5 day; any GPU pilot follows only if overlap is demonstrated. No overlap means no dose-independent causal claim.