Response to Reviewers — 1008 second-round revision

Round framing. The manuscript designated 01549 was submitted on 7 September 2026 after the first review round. The R1/R2/R3 comments addressed here concern the subsequent second round. We retain the prior 01549 work that remains relevant, including the preference study, FAC re-examination, original 276-object analysis, and historical adapter results; none is presented as a new Fresh C result.

> **Author verification required before submission.** The retained `CAG-D-26-00962-reviews.pdf` is labeled as initial-submission reviews, and no exact second-round export or editor email is available in this workspace. The round mapping is author-confirmed in the retained issue map; the headings below paraphrase that map and are not represented as verbatim reviewer quotations. Compare them with the editor's original second-round decision before finalizing this response.

Dear Editor and Reviewers,

We thank you for the careful assessment. This revision keeps Timestep-Conditioned Adapter Scaling (TCAS) as the method and narrows its claims to the evidence available. We now describe the actual Fresh C figures and controls, correct the stage-placement result, distinguish applied scale from schedule shape, and replace an invalid interaction-bootstrap inference with the available exploratory object-cluster robust Wald analysis. The revisions and remaining limits are detailed below.

Reviewer 1

Comment 1 — Visual artifacts, full-object comparisons, and 3D fidelity

Response. We agree that the earlier panels did not make it easy to compare the complete object and the same local detail. Figure 7 now shows reference, GFL, GFH, and C3 for a C3–GFL endpoint-favorable case and an adverse case, with each full target-view context paired with the same pixel ROI. Figure 4 preserves the fixed-scale comparison and adds a separate Generic Linear counterexample; its RGB outputs are not independently adjusted. Figure 6 reproduces the historical pink/purple failure from the submitted 01549 Figure 4, but its original panel-level object/run mapping is unavailable, so we use it only to document the historical failure. Supplementary Figure S1 shows two selected examples across six views and seven conditions, including No Adapter, LLH, and Generic Linear. These outcome-selected examples do not estimate artifact prevalence or texture fidelity.

The unmodified native-default adapter path is not represented by Fresh C's

No Adapter condition: No Adapter sets scale to zero, while an unset scale field applies the geometry residual without a scale multiplier. Fresh C has no separate output for that native-default path, so this comparison remains unresolved. The historical N=20 unseen-view check uses unlit base color and has mixed endpoint directions. The attempted Fresh C bake did not pass its native-input gate; no valid Fresh C bake, seam, or PBR result is claimed. We also cannot verify whether the reviewer-facing portal exposed the supplied video.

A later four-object color-pathway development audit further distinguishes a control channel from a repair. Perturbing the global appearance embedding shifted chroma in the registered directions on all four selected objects. The Fresh B input-only audit found that stored embeddings matched neither compared source-conditioning reconstruction within 0.03 for 150/150 inputs, but did not establish that this discrepancy causes the artifacts. The C1 chroma anchor met its numeric gate (mean unseen-view ΔFG-CIEDE2000 -1.5602; 95% CI $[-2.2621, -0.6299]$) but failed locked visual review: one case showed no clear gain, while the other retained mismatch and showed a faint cyan/blue rim. B2's selected raw-source embedding substitution also failed its numeric gate (mean object-level change $+0.1729$; 95% CI $[-0.2044, +0.6810]$; one of four improved) and showed no clear visual correction. No independent Fresh B repair validation was performed. These development findings identify a controllable appearance pathway and a plausible cache-alignment contributor, not a validated color fix or complete root cause.

We also completed a separate, failure-enriched 26-object color diagnosis using paired No Adapter, GFL, and LLH outputs with six unique target views per object. GFL reduced mean reference-relative foreground CIEDE2000 versus No Adapter by 19.65 (95% paired object-bootstrap CI $[-26.37, -13.25]$; 24/26 objects), but the two cases shown in the submitted 01549 Figure 4 retained visible pink or purple hue shifts across all three conditions. LLH minus GFL increased mean CIEDE2000 by 3.58 (95% CI $[+0.46, +6.71]$) and reduced local gradient and high-frequency measures, consistent with detail suppression rather than a color correction. A decoded RGB snapshot showed the hue shift before baking; a ground-truth VAE round trip did not reproduce the same direction or magnitude in the selected case. Because this diagnostic subset was failure-enriched, it is not an independent population confirmation. The complete color-generation cause remains unidentified, and no repair is validated.

Location. Main manuscript: Figure 4 (p. 13), the archival color-shift Figure 6 (p. 15), full-object Figure 7 (p. 16), the color results and Table 7 (pp. 8–9), and Limitations (p. 11). Supplement: selected six-view/seven-condition examples in Figure S1 (p. 10), Sections S6–S8 (pp. 9–10), and the paired color diagnosis in S13 (p. 13). The Fresh C figure-source and ROI manifest records image hashes and crop coordinates.

Comment 2 — Held-out matched evaluation

Response. The earlier 01549 revision reported a probe-disjoint 276-object schedule-transfer analysis using a separate v2 strong-residual, per-layer-capped adapter. We retain that first-round result in Supplementary Section S9; it does not substitute for a matched comparison of the current main adapter.

Fresh C supplies a matched held-out image-space comparison on 300 UIDs disjoint from the earlier 24-object schedule-development probe and the original 300-object evaluation. Its primary four-condition run and same-cohort Generic Linear/LLH addendum share a frozen checkpoint, runner source, cohort, inputs, views, metric implementation, and per-object seed policy. The integrity gates pass for 1,200/1,200 primary rows and 600/600 addendum rows. The registered C3−GFL FG-PSNR difference is +0.522352 dB (95% object-bootstrap CI [+0.404160,+0.640669]).

We report the counter-results alongside this positive contrast: GFH and Generic Linear have more favorable cohort means than C3 on all four displayed endpoints. We also clarify that the comparison does not isolate schedule timing. The 50-step applied deep and middle scale sums are 82.5 each for C3 and 62.5 each for GFL; both have a shallow sum of 40 after the native 0.8 cap. Fresh C therefore supports the recorded C3 intervention relative to GFL under this wrapper, not a timing-only effect.

The Fresh C checkpoint is a controlled retraining; it is not presented as the unavailable historical 01549 main-table weight artifact or as a numerical reproduction of the submitted +0.96 dB estimate. The historical v1 main-table and v2 Table 9/FAC adapter identities remain separate. Core-7 is retained only as a retrospective Supplement diagnostic: it reports C3−GFH +0.725229 dB Full-view PSNR and −1.214235 dB FG-PSNR, but does not establish one shared runner or complete input identity. Fresh B and A3b results are also retained as sensitivity evidence rather than preregistered confirmation because the lock text's stated freeze time (2026-10-05 08:40 UTC) predates the earliest preserved commit containing that file (2026-10-06 00:52 UTC), so the retained history does not authenticate pre-run availability.

Location. Main manuscript: Methods (pp. 3–5), Tables 3–4 (p. 8), temporal-placement analysis (pp. 7–9), and Limitations (p. 11). Supplement: S2 (p. 2), S3 (pp. 3–5), S4–S5 (pp. 6–8), S8 (p. 10), and S9 (p. 11).

Comment 3 — Texture variation versus texture fidelity

Response. We agree that larger gradients, RGB variation, or Laplacian variance do not establish faithful texture reproduction. The manuscript treats these as variation diagnostics only and does not use them to claim fidelity. Fresh C reports image-space agreement endpoints, not a validated texture-fidelity measure. A separate retrospective analysis on FRESH_CONFIRM_300 found a 9/10 plus-one-corrected Holm association family between LLH−GFL image-space

deltas and ground-truth descriptors; its Laplacian-variance strata range from a +4.14 dB FG-PSNR mean (92% LLH-favorable) in Q1 to slightly negative means in Q3 and Q4. The only non-rejecting test was FG-LPIPS versus foreground coverage (Holm-adjusted p=0.0508). This is a noncausal, post-outcome sensitivity result, not a fidelity measure, a selector, or evidence for C3. The earlier 3AFC study remains a separate subjective-preference result for its original three conditions; it is not Fresh C evidence, a reference-fidelity measure, or a 3D-quality result.

Location. Main manuscript: metric definitions (p. 3), Figures 4–7 (pp. 13–16), preference Table 9 (p. 10), and Limitations (p. 11). Supplement: S2 (p. 2), S7 (pp. 9–10), and the retrospective complexity analysis in S12 (p. 12).

Comment 4 — Paired uncertainty and Edge-SSIM

Response. We agree that an interval crossing zero does not establish equivalence or non-inferiority. Table 4 reports paired Fresh C intervals. C3−GFH Edge-SSIM is −0.01881 (95% object-bootstrap CI [−0.02084,−0.01683]), which we describe as an endpoint trade-off. No non-inferiority margin was prespecified, and we make no equivalence or non-inferiority claim. The historical 0.529-versus-0.559 values are not presented as a replicated Fresh C comparison.

Location. Main manuscript: Table 4 and endpoint interpretation (p. 8). Supplement: Section S2 (p. 2).

Comment 5 — FAC reproducibility and code availability

Response. We retain the negative FAC controlled re-examination reported in the first-round 01549 revision, including its seven-configuration results and untrained control. The stated study used roughly 6×10^3 trainable controller parameters, AdamW with a peak learning rate of 10^{-4}, 100 warm-up steps, 2,000 training steps, a reported 1,706-object training pool, and a paired 300-object evaluation. The reported trained configurations were 1.10–2.94 dB below TCAS on FG-PSNR; the untrained 12-object C3-envelope control had ΔFG-PSNR −0.18 dB (95% CI [−0.42,+0.05]).

We disclose that the recovered configuration points to a 1,200-object list rather than the 1,706-object pool stated in 01549, and the available archive does not bind the training seed, outputs, and runtime in a single run manifest. These remain previously reported results, not a new 1008 experiment. The earlier reviewer encouraged code release; we do not claim a public release.

Location. Main manuscript: FAC discussion (p. 10) and Limitations (p. 11). Supplement: Section S11 and Table S13 (pp. 11–12).

Reviewer 2

Concern 1 — Methodological contribution beyond schedule selection

Response. The contribution is a task-specific study of how geometry-residual strength can be controlled and measured inside a multi-view texturing system. TCAS makes the requested intervention explicit, while the execution analysis distinguishes requested scale, wrapper-applied scale under native caps, and the resulting residual magnitude. Fresh C then compares that intervention with fixed and scheduled alternatives under matched inputs and object-paired image-space endpoints. This combination provides reproducible evidence about what the control changes in this pipeline and about where interpretation fails: GFH and Generic Linear have more favorable means than C3 on all four displayed cohort endpoints, and C3/GFL differ in applied deep- and middle-layer scale sums, so their contrast does not isolate timing alone. We therefore retain TCAS as the named task-specific intervention and empirical comparison framework, without claiming that piecewise scheduling itself is a new general algorithm or that C3 is optimal.

We now compare the scope directly with Scheduled Style Injection (SSI), which studies style-injection schedules across layers and denoising time and also schedules geometry-conditioned ControlNet strength. SSI establishes that schedule design is not new in general. TCAS addresses a different task and residual interface. Its contribution is the operational specification of this intervention, the paired comparison of fixed and scheduled controls, and the resulting account of cap-, dose-, endpoint-, and object-dependent behavior in multi-view texture generation. The stage-placement analysis further shows that temporal L–L–H is favored over C3 on all six formal endpoints in a separate post-inspection follow-up. Fresh B's lock chronology is not authenticated, so those temporal and A3b findings are presented as sensitivity evidence; the retrospective texture-complexity analysis likewise identifies boundaries rather than a schedule selector. These results support an empirical system contribution, not a universal schedule law or a claim that C3 is the overall winner.

Location. Main manuscript: Introduction and Related Work (pp. 1–2), TCAS Method (pp. 3–4), Tables 3–8 (pp. 8–10), stage analysis (pp. 7–9), and Limitations/Conclusion (p. 11). Supplement: S3 (pp. 3–5), S4–S5 (pp. 6–8), S9 (p. 11), and the retrospective complexity boundary in S12 (p. 12).

Concern 2 — CAI and empirical schedule selection

Response. We agree that the CAI proposition and proof did not independently derive C3. We remove CAI's selector, proposition, proof, and optimality claim. We retain the TCAS residual-injection equation, three-stage profile, C3 definition, and 24-object development record, now identified as historical development evidence rather than a verified ranking or unique optimum. This correction concerns CAI's rationale and does not remove TCAS.

Location. Main manuscript: TCAS Method (pp. 3–4), development subsection/Table 2 (p. 7), and Figure 5 (p. 14). Supplement: historical development diagnostics in S10/Table S12 (p. 11).

Concern 3 — Cross-backbone generalization and interaction inference

Response. We agree that adapter variations within one MVPainter-style family do not establish cross-backbone generalization. The MV-Adapter probe analyzed 98 of 99 planned inputs; the omitted frozen GLB had path geometry but no triangle faces, and no substitute was used. The production additive-residual bootstrap had an invalid null for interaction inference and is not used as evidence.

The post hoc exploratory object-cluster robust Wald reanalysis (8 degrees of freedom) found no interaction surviving Holm correction across four metrics: FG-LPIPS raw p=0.101, Holm p=0.303; PSNR p=0.434, Holm p=0.868; FG-SSIM p=0.039, Holm p=0.156; and CIEDE2000 p=0.862, Holm p=0.862. These results are not proof that interaction is absent or that conditions are equivalent. A separate 76-object MV-Adapter layer-mapping study found mixed metric directions under a budget-neutral B' partition; its M3 companion changed both grouping and total nominal residual budget (mean multiplier 0.8065) and was classified as retrospective because the protocol and results first appear in the same preserved commit. The two mapping analyses do not establish a general mapping law or generalize TCAS across backbones. The separate MVDiffusion probe covered 75 objects through a different correspondence-aware interface and produced mixed or negative endpoint results. Because these probes use different conditioning paths and runners, they are interface checks rather than a matched architecture comparison; cross-backbone generalization remains unestablished.

Location. Main manuscript: Cross-Interface Evaluation/Table 8 (p. 10) and Limitations (p. 11). Supplement: Section S5 (pp. 6–8).

Reviewer 3

We thank Reviewer 3 for the positive assessment. No new technical change or experiment was requested.

Computers & Graphics (2026)

Contents lists available at ScienceDirect

Computers & Graphics

journal homepage: www.elsevier.com/locate/cag

Timestep-Conditioned Adapter Scaling for Multi-view Diffusion Texture Generation

ARTICLE INFO

Keywords: multi-view diffusion, 3D texturing, geometry adapter, residual scaling, timestep conditioning

ABSTRACT

Geometry-conditioned multi-view texturing uses adapter residuals to inject shape information into a denoiser, but a requested scale does not by itself specify the value applied by a wrapper or the resulting correction magnitude. We study Timestep-Conditioned Adapter Scaling (TCAS) on the existing GeoTex-Adapter path, distinguishing requested scale, native-cap applied scale, and post-scale residual magnitude. On the matched Fresh C cohort of 300 held-out objects, the recorded C3 profile $(1.25, 2.50, 1.25)$ exceeds global fixed-low scaling (GFL) on the registered FG-PSNR endpoint by +0.522352 dB (95% object-bootstrap CI [+0.404160, +0.640669]). Global fixed-high scaling (GFH) and Generic Linear have more favorable means than C3 on each of the four displayed endpoints. C3 and GFL also differ in applied deep- and middle-layer scale sums, so this contrast does not isolate temporal placement. In a separate post-inspection 276-object placement analysis, temporal L–L–H is favored on all six formal endpoints despite near-matched scalar sums; post-scale residual norms remain unmatched. The results characterize endpoint- and object-dependent responses of this implementation. Piecewise scheduling itself is established prior work, and the evidence does not establish an optimal schedule, a dose-independent effect, general transfer, or PBR/material fidelity. The contribution is a reproducible, graphics-task-specific account of how this control is requested, applied, and realized in image-space evaluation, with its downstream boundary made explicit.

1. Introduction

Multi-view diffusion systems generate target-view images that can later be projected or baked onto a 3D surface. In geometry-conditioned systems, an adapter injects shape information into the denoiser through intermediate residuals. The strength applied to those residuals is therefore a practical inference-time control: changing it can alter how strongly geometry conditions the generated views and can also change their appearance. The useful setting depends on the adapter implementation and on the image endpoint being considered, so a single structural score cannot characterize the full effect.

This paper studies that control question in a fixed MVPainter-style multi-view texturing pipeline. Timestep-Conditioned Adapter Scaling (TCAS) applies a piecewise requested-scale profile to the existing GeoTex-Adapter residual path without training parameters or changing model weights. The formulation records three separate quantities: the scale requested for a denoising step and layer group, the scale actually applied after native wrapper caps, and the norm of the scaled residual correction. This makes the implementation and the scope of each comparison measurable.

Layer- and timestep-dependent control is established prior work. Scheduled Style Injection (SSI) varies StyleID's attention-mixing parameter across decoder layers and denoising steps, and also schedules geometry-conditioned ControlNet strength along those axes. We therefore do not present TCAS as a new general scheduling primitive. Our contribution is an empirical graphics-system characterization of a specific residual path: we measure how requests pass through native caps, how the applied corrections vary, and how fixed and scheduled controls compare on paired multi-view image endpoints. SSI's style-transfer results do not determine the results of these object-level texturing comparisons; differences in task and interface require measurement here, but do not make the underly-

ing scheduling idea new.

We evaluate the recorded C3 profile against fixed and alternative scheduled controls on the held-out Fresh C cohort. The registered C3–GFL FG-PSNR result is reported with the stronger GFH and Generic Linear mean results and with the applied-scale difference that prevents a timing-only interpretation. A separate post-inspection placement analysis favors temporal L–L–H on six endpoints, while layer-window maps and cross-interface probes show that responses depend on the tested intervention and interface. These results support an empirical system contribution, not a dose-independent schedule law. Fresh C measures target-view images; it does not establish final baked appearance or PBR/material fidelity.

The contributions are:

(1) **Control formulation.** We define and implement three-stage, timestep-conditioned scaling on an existing geometry-adapter residual path, explicitly distinguishing requested scales, applied scales after native caps, and post-scale residual corrections. The Fresh C execution records and the 50-step, three-layer accounting in Supplementary Section S4 make these semantics reproducible for the tested wrapper; the formulation is an operational description, not a new general scheduling operator.

(2) **Controlled empirical characterization.** We provide a paired 300-object comparison of fixed and scheduled controls, including the registered C3–GFL FG-PSNR result, the stronger GFH and Generic Linear means, and the LLH endpoint trade-off (Table 4). We separately report strict-276 temporal placement and layer-window sensitivity, identifying where nominal applied-scale sums or realized residual magnitudes prevent a timing-only or dose-independent interpretation (Table 5; Supplementary Sections S3–S5).

(3) **Graphics-oriented validation and limits.** We distinguish image-space agreement, reference-color agreement, texture variation, subjective preference, and baked 3D appearance. Paired full-object/same-ROI examples, the failure-enriched color diagnosis, and interface-specific adapter comparisons document appearance failures and object-dependent responses (Figures 6 and 7, Table 7, and Table 8). They also show why target-view image improvements cannot be read as proof of baked 3D or PBR/material fidelity.

2. Related Work

2.1. 3D Texture Generation and Multi-view Diffusion

Single-image 3D generation is often decomposed into geometry generation and texture generation. Single-image-to-3D, multi-view diffusion, and sparse-view reconstruction methods have improved structural stability and view synthesis [1, 2, 3, 4, 5, 6, 7, 8]. Texture synthesis remains a distinct challenge for asset realism and usability [15]. Many systems generate RGB views with 2D diffusion and project or bake them onto a surface [9, 10, 11, 12]. Recent work also explores exemplar-based texture synthesis and diffusion-based procedural materials [16, 17].

Paint3D, MVPaint, Hunyuan3D 2.0, and MVPainter use geometry-conditioned multi-view generation in 3D texturing pipelines [18, 19, 20, 13]. MV-Adapter, MVDiffusion, and other synchronized multi-view methods address consistency through different adapter and denoising designs [21, 22, 23, 24]. These systems establish the graphics context for this work. The present evidence does not compare final baked-texture quality across those systems.

2.2. Adapter-Augmented Diffusion Models and Geometric Control

T2I-Adapter and ControlNet inject spatial conditions into pretrained diffusion models through controllable branches [25, 26]. IP-Adapter introduces image prompts through decoupled attention [27]; LoRA and related methods use scaled low-rank updates [28, 29, 30]. These methods establish condition and residual scaling as common implementation controls. TCAS does not introduce adapter conditioning or a new multi-view backbone; it studies a requested scale on the existing GeoTex-Adapter residual path.

2.3. Diffusion Guidance Scheduling and Timestep Control

Diffusion denoising is often discussed in terms of coarse-to-fine evolution, and prior work varies guidance strength over sampling progress to study quality, diversity, or fidelity [32, 33, 34, 35]. Limited-interval classifier-free guidance is one example; it changes the combination of conditional and unconditional score predictions [14].

SSI is the closest scheduling precedent. It varies StyleID's query-mixing parameter across decoder layers and denoising time, tests schedule direction and shape, and applies scheduling to a separate geometry-conditioned ControlNet branch; its reported ControlNet scheduling gains over fixed conditioning are modest relative to its style-injection results [31]. This prior work already covers the general idea of layer/time scheduling and geometry-conditioning schedules. TCAS acts on the residual correction produced by GeoTex-Adapter from normal, depth, and foreground-mask conditions, with group-specific wrapper caps; it does not modulate StyleID attention mixing or add a ControlNet branch. The task, interface, and measured endpoints differ, so SSI does not imply the Fresh C paired estimates, the requested/applied/residual accounting, or the observed image- and color-agreement failures. These task-specific measurements are the empirical contribution; the scheduling concept itself is not claimed as new.

2.4. Multi-view Adapter Interfaces

MV-Adapter uses a condition guider for camera or geometry maps together with decoupled spatial, multi-view, and optional image cross-attention layers [21]. MVDiffusion uses correspondence-aware attention modules to connect pixels across views [22]. These interfaces differ from the GeoTex-Adapter residual branch studied here. We therefore report each auxiliary intervention within its own tested implementation and cohort, without pooling metrics or interpreting it as an architecture-matched replication of TCAS. The main cross-interface results are summarized in Table 8.

3. Method

3.1. Task Definition and Base Pipeline

This paper builds on the MVPainter-style geometry-controlled multi-view texture generation framework [13] and studies how the inference-time strength of geometric condition injection affects the generated results. Figure 1 illustrates the overall pipeline: the adapter injects geometric residuals into UNet features, TCAS adjusts their strength across denoising stages, and the final multi-view outputs are projected and baked into a textured 3D mesh. Given a single reference image and the corresponding 3D geometry, the base pipeline first generates RGB texture images from six target views and then obtains a textured mesh through projection or texture baking. We do not modify the base UNet, VAE, or pretrained adapter weights. Instead, the adapter scaling factor is treated as the key inference-time control variable.

During multi-view diffusion, the 3D geometry is rendered into normal maps, depth maps, and foreground masks, denoted as:

$$\mathbf{G} = \{G_{\text{normal}}, G_{\text{depth}}, G_{\text{mask}}\}. \tag{1}$$

The geometric adapter maps these conditions to a residual correction term and injects it into the intermediate hidden states of the UNet. For the t-th denoising step, the adapter injection can be written as:

$$\mathbf{h}'_{t,\ell} = \mathbf{h}_{t,\ell} + s_{\text{app}}(p_t, \ell)\, A^{\ell}_{\phi}(\mathbf{h}_{t,\ell}, \mathbf{G}), \tag{2}$$

Here ℓ indexes an injected layer group and p_j is normalized sampling progress at inference step j. The *requested scale* $s_{\text{req}}(p_j, \ell)$ is the value assigned by the profile; the *applied scale* is the value used by the wrapper,

$$s_{\text{app}}(p_j, \ell) = \min\big(s_{\text{req}}(p_j, \ell), c_\ell\big), \tag{3}$$

when the native cap c_ℓ is active, and otherwise equals the request. The resulting correction is $\mathbf{r}_{j,\ell} = s_{\text{app}}(p_j, \ell) A^{\ell}_{\phi}(\mathbf{h}_{j,\ell}, \mathbf{G})$; its *post-scale residual magnitude* is $m_{j,\ell} = \|\mathbf{r}_{j,\ell}\|_2$. Thus neither a requested nor an applied scale label alone determines the correction magnitude.

In the audited Fresh C wrapper, the three cap groups map the GeoTex residual wrappers on the UNet ResNet blocks as follows: deep comprises the UNet middle block and `up_0`, middle is `up_1`, and shallow is `up_2`. These names identify the implementation's grouped injection points; they do not assert universal semantic roles for denoising stages. The corresponding Fresh C caps are 3.0, 3.5, and 0.8. C3's shallow request of 1.25 is therefore applied as 0.8, while its deep and middle requests pass unchanged. GFL and GFH also clip at the shallow cap. These cap values describe the audited Fresh C wrapper; row-level cap semantics for the historical 24-object probe are not fully recoverable. TCAS scales a geometry residual added to feature activations; it does not reweight conditional and unconditional score predictions as classifier-free guidance does.

3.2. Image-Space Metrics and Texture-Variation Diagnostics

We use separate terms for separate outcomes. *Image-space agreement* refers to target-view PSNR, LPIPS, and SSIM measurements on the named image regions; it is not a measure of final 3D material quality. *Reference-color agreement* refers to color differences such as CIEDE2000 computed against the corresponding reference. *Texture variation* refers to statistics such as RGB standard deviation, Laplacian variance, gradient magnitude, and high-frequency energy; these may rise with detail or artifacts and are not fidelity measures. *Subjective preference* refers only to the submitted forced-choice study. *Baked 3D appearance* refers to rendered mesh outputs after texture projection or baking, while *PBR/material fidelity* additionally concerns appearance under the material and illumination model. These endpoints are not interchangeable. FG-SSIM, Edge-SSIM, and Crop-SSIM quantify structural similarity on the indicated image regions; PSNR measures pixel error; LPIPS compares learned image features [36, 37].

Laplacian variance, RGB standard deviation, and gradient magnitude can describe local high-frequency or color variation. For example, a foreground Laplacian statistic may be written as

$$M_{\text{var}} = \text{Var}(\text{Lap}(I_{\text{out}} \odot G_{\text{mask}})), \tag{4}$$

where G_{mask} is the foreground mask. A larger value is not necessarily better: detail, noise, repeated patterns, or artifacts can all increase local variation. We therefore do not interpret these proxies as texture fidelity or use them alone to claim that a schedule better preserves reference appearance.

3.3. TCAS Control Semantics and Identifiability

TCAS replaces a globally fixed requested adapter scale with a piecewise profile over normalized sampling progress. The early/middle/late partition is an operational description, not a universal claim about the function of every denoising stage. C3's low–high–low profile is the specific development choice evaluated in this revision; it is not presented as an independently derived optimum.

To avoid ambiguity caused by different raw timestep directions in different schedulers, $p_j = j/(T-1)$ denotes execution progress for inference step $j \in \{0, \ldots, T-1\}$. Thus $p = 0$ is the first update from the initial noisy latent and $p = 1$ is the last denoising update before decoding. For $T = 50$, TCAS divides execution into steps 0–16, 17–32, and 33–49 (17/16/17 steps):

$$s_{\text{req}}(p) = \begin{cases} s_e, & 0 \le p < 1/3 \\ s_m, & 1/3 \le p < 2/3 \\ s_l, & 2/3 \le p \le 1 \end{cases}. \tag{5}$$

We call this variation across the 50 inference steps the *temporal allocation*. A *layer allocation* varies the requested scale across the three wrapper groups at a fixed step. For a schedule with applied values $s_{\text{app}}(j, \ell)$, its per-group applied-scale sum is $B_\ell = \sum_j s_{\text{app}}(j, \ell)$; the unweighted sum $B_{\text{nom}} = \sum_\ell B_\ell$ is a *nominal applied-scale budget*. It does not weight scales by the adapter activations. The logged vector of post-scale magnitudes $\{m_{j,\ell}\}$, or a stated aggregation of that vector, is the *realized residual dose*. It depends on the current hidden state and on earlier denoising updates, so equal nominal budgets do not imply equal realized doses.

The selected C3 configuration used in this paper is:

$$(s_e, s_m, s_l) = (1.25, 2.50, 1.25). \tag{6}$$

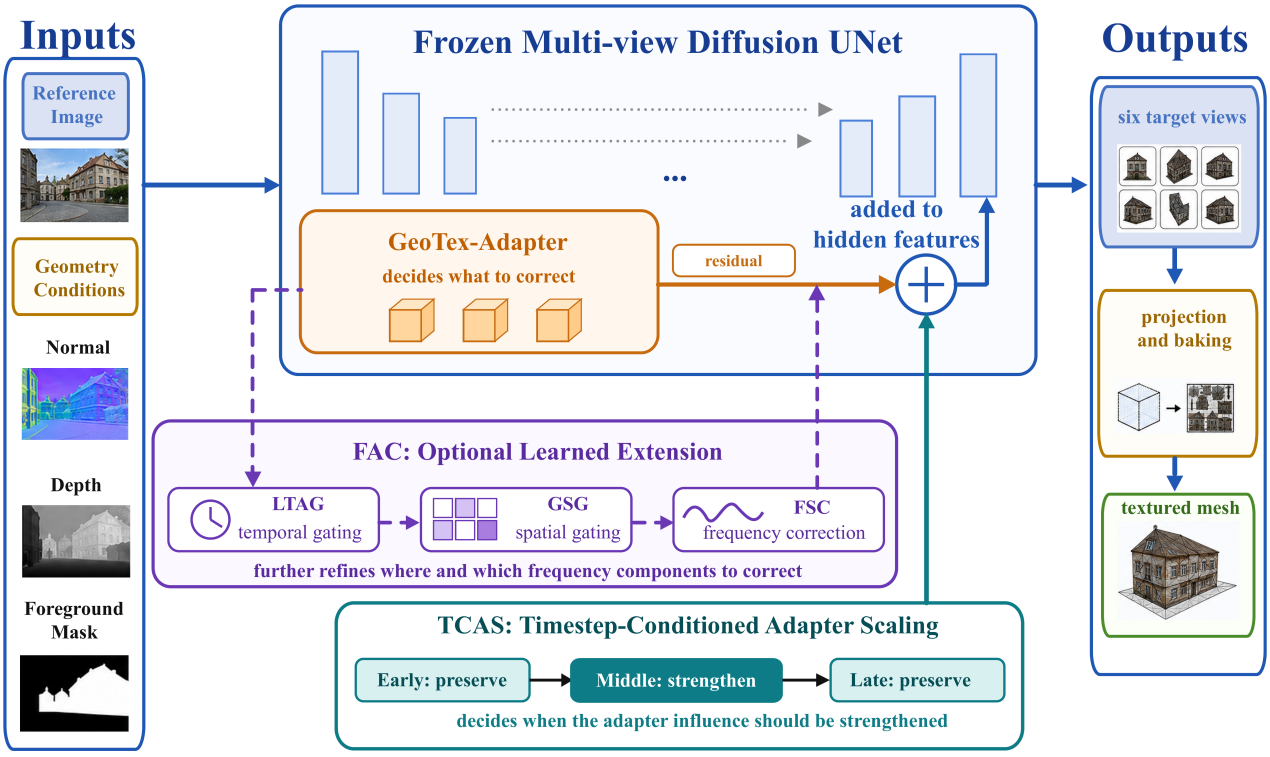


Fig. 1. Method overview reproduced from the authenticated 01549 submission. The optional FAC branch is retained as a previously reported extension; its controlled negative re-examination from 01549 is summarized in Section 4.10. TCAS remains the paper's training-free method. Projection and baking are pipeline context, not outcomes validated by Fresh C image-space evidence.

The low–high–low profile is a concrete inference-time intervention motivated by the submitted development analysis. Its stage interpretation is a task-specific hypothesis: the tested early, middle, and late allocations may produce different responses. Fresh C evaluates the recorded profile against matched controls but does not prove a universal function for any stage. TCAS adds no trained parameters and does not modify base-model or adapter weights; it replaces the requested scale during inference.

What C3 changes and what its comparisons identify. C3 requests (1.25, 2.50, 1.25) for each of the 50 steps across the three layer groups before native caps. In Fresh C, the applied deep and middle profiles retain those values; the shallow profile is clipped to 0.8 throughout. GFL requests 1.25 at every step and group, so its deep and middle values are 1.25 throughout and its shallow value is also clipped to 0.8. Their per-group applied-scale sums are therefore 82.5 versus 62.5 in deep and middle, and 40 for both in shallow. The C3–GFL contrast changes the middle-third deep/middle scale magnitude as well as its temporal allocation; it is not a timing-only comparison.

To identify a temporal-placement contrast, controls must hold the per-layer applied values and nominal per-group sums fixed while moving the same values between time windows. To identify a scale-strength contrast, controls must hold the layer and temporal allocations fixed while changing amplitude. Both contrasts also require reporting post-scale residual magnitudes: matching requested values or applied-scale sums does not match the evolving residual corrections. The strict-placement analysis approaches scalar-sum matching but does not match realized residual dose; post-hoc adjustment for that state-dependent quantity does not by itself recover a dose-independent causal effect.

Development origin and scope of C3. The submitted 01549 paper records C3 = (1.25, 2.50, 1.25) as the selected profile from a 24-object development probe containing 17 uniform, layer-wise, and time-conditioned variants. We retain that development record because it explains the origin of the profile studied here. The complete row-level outputs and selection-time provenance are not available in the current evidence authority, so the historical ranking and labels cannot be independently checked. We therefore make no formal optimality or uniqueness claim based on this probe. Fresh C evaluates the recorded profile without further schedule search.

3.4. Geometry Adapter and Inference Interface

GeoTex-Adapter is treated as the fixed geometry-residual backbone, not as a new contribution of this paper. It receives a five-channel input comprising a three-channel normal map, a one-channel depth map, and a one-channel foreground mask. A multi-scale encoder produces residual features that are injected into the frozen UNet. TCAS acts on the requested multiplier along denoising progress; it does not retrain the adapter, the UNet, or the VAE and introduces no learned parameters.

The first-round 01549 revision also reported a learned Full Adaptive Correction (FAC) extension. We retain its controlled negative re-examination as a previously reported auxiliary result in Section 4.10 and Supplementary Table S11; it is not a new 1008 training run. TCAS remains the paper's sole training-free method.

4. Experiments

4.1. Experimental Setup

4.1.1. Experimental protocol and comparisons

The original 01549 submission reports fixed-scale experiments, a 24-object schedule-development probe, and a 300-object pooled evaluation. The source figures and tables are authentic submission artifacts, but the current evidence archive does not contain a complete object-level manifest and run record for every historical result. We therefore preserve the fixed-scale and C3-selection questions as historical development evidence and do not present their unrecovered numbers as newly verified confirmatory results.

Fresh C is a held-out, matched image-space study on 300 frozen object identifiers, independent of the original 24-object schedule-selection probe and the original 300-object evaluation cohort. The primary four-condition run (No Adapter, GFL, GFH, and C3) and the same-cohort Generic Linear/LLH addendum use the same frozen adapter checkpoint (SHA-256 `0618d6b284ab47aa16d7a89dc447f5ba4455ff6c9d918e5e129478a99e0114c0`), runner source (SHA-256 `e5c28e915ba137d3541eb4c032daaf62a798332d63b32917b749fe1bcc74e9e3`), 300-object cohort, reference/geometry inputs, target views, metric implementation, and per-object seed policy 42 + `object_idx`. The verified gates cover 1,200/1,200 primary rows and 600/600 addendum rows. The registered primary endpoint is paired C3–GFL FG-PSNR. Fresh C's 300 UIDs have zero overlap with the original evaluation cohort (including its probe), the original 1,118-object training list, and the clean-v2 evaluation cohort. The Fresh C comparison uses one frozen checkpoint and common execution protocol; hash identity with the historical Table 4 checkpoint is not used as a validity condition. The historical distinction remains: 01549 names a v1 checkpoint for its main tables and a separate v2 strong-residual, per-layer-capped adapter for Section 4.7/FAC; Fresh C is a controlled retraining, and its estimate is not presented as a numerical reproduction of the old +0.96 dB contrast. The separate post-inspection strict-276 analyses use the clean-v2 object list and record the same checkpoint SHA as Fresh C; their cohort and run provenance differ from Fresh C, and they are reported only as bounded supporting diagnostics.

4.1.2. Evaluation sets and metrics

The submitted paper describes a 50-object, 13-scale sweep; a purposive 26-object texture audit; a 24-object, 17-variant development probe; and a pooled 300-object evaluation. We retain the 50-object sweep and the 24-object probe as historical diagnostics because they motivate the scale-allocation question. Their object-level raw outputs, full selection provenance, or metric implementation records are incomplete in the current archive. The 26-object audit's object IDs, selection trace, and proxy definitions also cannot be reconstructed; its prevalence and texture-loss claims are withdrawn.

Fresh C adds a revision-era 300-object image-space comparison with foreground PSNR, LPIPS, SSIM, and Edge-SSIM. These measure agreement with target-view renderings on their named endpoints; they do not measure texture fidelity, material appearance, seams, or human preference. The separately reported 01549 preference study is a subjective choice task, not a texture-fidelity measure. Laplacian variance, RGB standard deviation, and gradient magnitude are treated only as variation diagnostics, since larger values can also reflect artifacts. Object is the inferential unit; views are averaged within object and are not independent samples.

4.1.3. Implementation and dataset details

The authenticated 01549 source names `geotex_refattn_v1/geotex_step_0002000.pt` for its main tables but provides no checkpoint hash. Fresh C uses a controlled-retraining GeoTex-Adapter checkpoint with SHA-256 `0618d6b284ab47aa16d7a89dc447f5ba4455ff6c9d918e5e129478a99e0114c0`, trained from the same 1,118-object UID list (SHA-256 `7b3e998fc494ecffde247793e24efab715b4536b96382eb5fe559434fa57cae7`). The primary run and same-cohort addendum share this frozen checkpoint, runner source, seed policy, inputs, and endpoint implementation. Fresh C uses the MVPainter-style joint six-view latent UNet and VAE, EulerDiscreteScheduler for 50 steps, six fixed target views at 256×256, and per-object seed 42 + `object_idx`, shared across all compared conditions for that object. The 300-object manifest, per-condition row counts, and input/output hashes are indexed in the Supplementary reproduction map.

Fresh C's native wrapper caps are deep 3.0, middle 3.5, and shallow 0.8. The recorded scale request, applied value, and post-scale residual norm are separate fields. The requested C3 profile is $(1.25, 2.50, 1.25)$; its shallow value is clipped to 0.8 in this run. The Generic Linear and LLH addendum conditions are separately defined and do not equalize realized residual dose. Detailed step- and layer-group records are supplied in the Supplement.

FG-PSNR, FG-LPIPS, FG-SSIM, and Edge-SSIM are computed per object after averaging the six target views. The registered primary comparison is C3–GFL FG-PSNR; secondary and addendum tests follow their locked comparison families as documented in Supplementary Section S2. Fresh C's primary endpoint is foreground-masked PSNR. The historical Table 4 implementation and the Fresh C endpoint are not claimed to be numerically identical; the Fresh C estimate is a new matched held-out result, not a reconstruction of the old +0.96 dB contrast. The original 50-object sweep and 24-object probe used separate historical records whose complete current row-level lineage is unavailable, so their implementation details are not inferred from Fresh C.

Statistical reporting. Fresh C reports the object-level paired mean difference and percentile-bootstrap 95% interval from

Table 1. Historical 01549 report of a 13-scale sweep on the stated 50-object set. Values are reproduced from the submission and have not been independently recomputed from row-level outputs; they are retained as development diagnostics.

Scale	ΔFG-SSIM↑	POS	ΔCrop-SSIM↑	ΔPSNR↑
0.75	−0.013	20/50	+0.113	7.53
0.90	+0.024	30/50	+0.137	8.32
1.00	+0.047	34/50	+0.147	8.79
1.10	+0.070	36/50	+0.159	9.08
1.15	+0.071	36/50	+0.157	9.25
1.25	+0.088	37/50	+0.161	9.37
1.35	+0.100	39/50	+0.163	9.63
1.50	+0.113	40/50	+0.163	9.76
1.60	+0.120	42/50	+0.166	9.83
1.75	+0.127	43/50	+0.173	9.97
2.00	+0.134	44/50	+0.169	9.96
2.25	+0.142	43/50	+0.175	9.91
2.50	+0.145	44/50	+0.165	9.91

10,000 object resamples. The registered primary C3–GFL FG-PSNR contrast is unadjusted as the single primary endpoint. Secondary and addendum families and their multiplicity adjustments are detailed in Supplementary Section S2. CIs outside that prespecified family are pointwise and descriptive. The earlier three-condition 3AFC study is reported separately below as a subjective preference result; it is not a Fresh C endpoint or evidence of texture fidelity.

4.2. Shape-Texture Trade-off of Adapter Scaling

Table 1 reproduces the 13-scale sweep reported in the submitted 01549 manuscript. We retain it to preserve the original development question, but the object-level outputs, full run manifest, and metric implementation mapping are not available in the current authority. Its values and Figures 2–3 are therefore historical diagnostics, not a new reanalysis or confirmatory evidence.

The submitted table reports a rising ΔFG-SSIM over most of the listed scale range and a near-plateau in the crop metric at higher settings. We retain this reported pattern only as historical motivation: absent row-level outputs and a complete run manifest, we do not treat it as an independently confirmed scale-response law. It motivates asking whether a single aggregate structural score fully describes the effect of adapter scaling.

The submitted heatmap depicts object-level variation across the reported scales, but its underlying rows cannot be reconstructed from the current evidence archive. We retain the visual diagnostic as part of the 01549 research history and do not infer a population-level response distribution from it. The updated Fresh C analysis below uses its own frozen cohort and matched inputs.

4.3. Timestep-Conditioned Adapter Scaling: Shape-Texture Balance of C3

The submitted paper records 17 uniform, layer-wise, and timestep-conditioned variants on a 24-object development probe and identifies C3 as its chosen profile. This study question remains relevant: does a stage-dependent request offer a useful control setting beyond a single global scale? Because the complete historical row-level data and selection provenance are not available in the current authority, the original Table 2 and Figure 5 are retained as historical development records, not as a verified ranking or proof of optimality. Fresh C evaluates the recorded C3 profile without another schedule search.

The original table and plot use shape and texture-variation proxies to describe the development choice. They do not establish that texture was faithful to the reference, nor that C3 was the unique best candidate. In this revision, C3 is therefore described as the historical profile carried forward for matched evaluation, not as a theorem-derived or globally optimal schedule.

The requested profile is piecewise constant over normalized

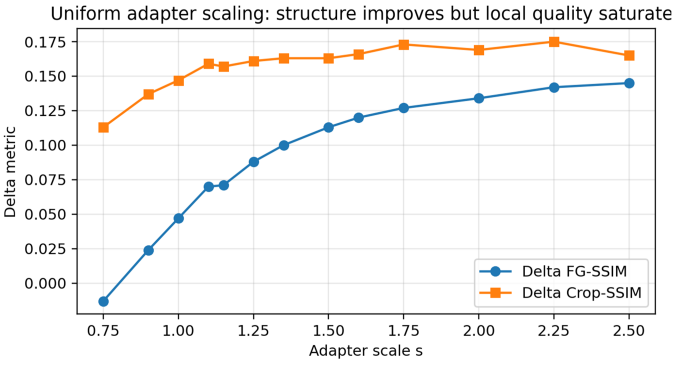


Fig. 2. Historical fixed-scale diagnostic from the submitted 01549 paper. The plotted source reports FG-SSIM and Crop-SSIM over adapter scales; the underlying object-level run records are not available for independent reanalysis. It motivates measuring more than one endpoint, but is not a current confirmatory result.

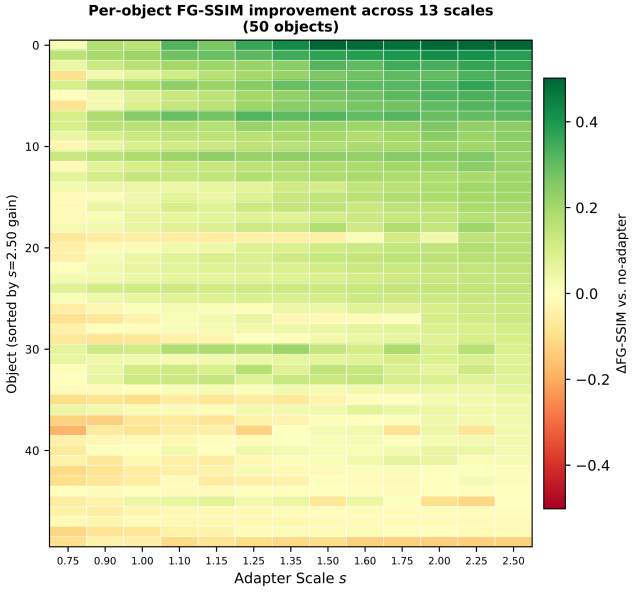


Fig. 3. Historical per-object ΔFG-SSIM heatmap reproduced from the 01549 submission (50 objects, 13 scales). The plot illustrates the originally reported response variation; row-level provenance is incomplete, so it is not used for a current population inference.

Table 2. Representative values reported for the historical 24-object, 17-variant development probe. The full candidate set is shown in Figure 5. The underlying row-level outputs and selection-time provenance are not available for independent reanalysis; the table records the submitted development result and does not establish an optimum.

Variant	Reported category	ΔFG-SSIM	ΔLap Var
C4	Late-stage high	0.105	−0.0120
C2	Timestep schedule	0.104	−0.0110
B2	Layer-wise schedule	0.090	−0.0135
A_s2.00	Uniform	0.087	−0.0062
C3	Mid-stage high	0.084	−0.0006
A_s2.25	Uniform	0.079	−0.0096
A_s2.50	Uniform	0.078	−0.0108
B6	Layer-wise schedule	0.000	−0.0027

sampling progress p:

$$s_{\mathrm{req}}(p) = \begin{cases} 1.25, & 0 \le p < 1/3, \\ 2.50, & 1/3 \le p < 2/3, \\ 1.25, & 2/3 \le p \le 1, \end{cases} \qquad \mathrm{C3} = (1.25, 2.50, 1.25). \tag{7}$$

This is an inference-time instruction on the existing adapter path; no weights are trained. Its tested value is task- and implementation-specific. The schedule operation itself is not claimed as a new general primitive.

Figure 4 presents compact Fresh C full-object and local-detail comparisons for the original C3-versus-fixed-scale question and the Generic Linear counter-result. The object-paired estimates in Section 4.4 carry the quantitative comparison.

The source Figure 5 (26-object texture ratios) is not included as an active result because its object IDs, selection trace, and metric lineage are incomplete. The original 24-object Table 3 is moved to Supplementary Section S10 as a verbatim historical diagnostic, with no texture-fidelity or population inference. Both source assets remain preserved in the authenticated submission and forensic archive; neither decision removes the fixed-scale research question or the development record of C3.

4.4. Large-Scale Matched Image-Space Evaluation

Fresh C is a revision-era matched image-space study on $N = 300$ objects. Its primary run contains No Adapter, global fixed-low (GFL), global fixed-high (GFH), and the historical C3 profile. A separate addendum evaluates Generic Linear and a layer-conditioned low–high–low profile (LLH) on the same 300 object identifiers; it is not an independent cohort. All conditions within a run share the object, reference/target inputs, six target views, initial latents, and seed. Object is the inferential unit, with views averaged within object.

The registered primary contrast is C3 minus GFL on foreground PSNR (FG-PSNR). The point estimate is +0.522352 dB (95% percentile-bootstrap CI [+0.404160, +0.640669]); C3 is higher on 206/300 objects (68.7%), with median paired difference +0.554514 dB. Registered secondary endpoints are reported with the multiplicity treatment and full paired statistics in Supplementary Section S2. PSNR, LPIPS, and SSIM quantify image-space agreement on their respective endpoints; they are not direct measures of texture fidelity or 3D material quality.

This primary contrast changes more than temporal placement. Across the 50 steps, the applied deep and middle scale sums are 82.5 each for C3 and 62.5 each for GFL; the shallow sum is 40 for both after the native 0.8 cap. The Fresh C estimate therefore supports the recorded C3 intervention relative to GFL under this wrapper, but cannot be attributed to timing alone or interpreted as a pure schedule-shape effect.

The C3–GFL comparison supports a bounded statement: on this frozen, UID-disjoint cohort and image-space endpoint, the recorded C3 intervention improves mean FG-PSNR relative to GFL. The paired conditions use one frozen checkpoint, common runner, seed policy, inputs, target views, metrics, and recorded native-cap semantics. The historical 01549 v1/v2 checkpoint distinction is preserved, and Fresh C is not presented as a numerical reproduction of the old Table 4 result. This historical difference does not alter the internal pairing or held-out status of Fresh C.

The counter-results delimit that finding. GFH has more favorable cohort means than C3 on all four displayed endpoints. Generic Linear also has more favorable means than C3 on all four; for C3 minus Generic Linear, the mean FG-PSNR difference is −0.593 dB (95% CI [−0.837, −0.351]), although the median is +0.073956 dB and C3 is higher on 155/300 objects (51.7%). LLH's FG-PSNR mean exceeds C3 by +0.621 dB, while its median difference is −0.153 dB and only 140/300 objects favour LLH; its lower mean FG-LPIPS favours LLH. For LLH versus Generic Linear, the FG-PSNR CI includes zero, which is not evidence of equivalence; the FG-LPIPS difference favours Generic Linear. The endpoint means and object-level distributions do not identify a universal winner.

Fresh C contains no validated texture-variation or texture-fidelity endpoint. We therefore do not revive the original 300-object RGB/Laplacian/gradient tables or their "retention" interpretation. In those diagnostics, a larger response could represent meaningful detail or an artifact; it cannot by itself establish fidelity.

4.5. Applied-Scale Budget and Temporal Placement

Fresh C compares the recorded C3 profile with GFL under the same frozen checkpoint, runner, seed policy, inputs, and target views. Its applied deep/middle scale sums are 82.5 for C3 and 62.5 for GFL, while the shallow sum is 40 for both after the 0.8 cap. The registered C3–GFL improvement therefore combines a change in middle-third strength with a larger per-group nominal budget; it does not isolate temporal placement.

A separate post-inspection placement analysis used the clean-v2 strict-276 cohort and a common runner within that run. It is not the Fresh C cohort, and it is not an independent replication of the historical 24-object selection probe. Temporal L–L–H means low requests during early and middle steps and high requests during late steps; it is distinct from the layer-conditioned LLH condition in the Fresh C addendum. With 50 steps, the strict-276 run assigns 17/16/17 steps to the three windows. C3 and temporal L–L–H have near-matched applied-scale sums (82.50 versus 83.75 in deep and middle; 40 in shallow), but their post-scale residual magnitudes are not matched.

Against the fixed-mean profile, C3's strict-276 FG-PSNR difference is +0.502875 dB (95% CI [+0.462018, +0.542244]).

Table 3. Fresh C image-space condition means. The primary run contains No Adapter, GFL, GFH, and C3. Generic Linear and LLH are revision-era addendum conditions evaluated on the same object identifiers, not a second independent cohort. Values are means of object-level metrics over six target views.

Method	N	FG-PSNR↑	FG-LPIPS↓	FG-SSIM↑	Edge-SSIM↑
Primary run					
No Adapter	300	6.9779	0.2288	0.4086	0.4314
GFL	300	11.3077	0.2051	0.3500	0.4728
GFH	300	12.5777	0.1891	0.4201	0.4961
C3 (TCAS)	300	11.8301	0.1998	0.3746	0.4773
Same-cohort revision-era addendum					
Generic Linear	300	12.4230	0.1862	0.4026	0.4880
LLH	300	12.4506	0.1876	0.3911	0.4863

Table 4. Fresh C paired differences, condition A minus B. Cells report the mean difference and pointwise 95% object-bootstrap CI. Positive PSNR/SSIM and negative LPIPS differences favour A. C3–No Adapter and C3–GFH are post-hoc contrasts; the addendum shares the primary run's object identifiers. Multiplicity-adjusted tests and the full endpoint family are in Supplementary Section S2.

Contrast	FG-PSNR (dB)	FG-LPIPS	FG-SSIM	Edge-SSIM
C3 − GFL	+0.522 [0.404, 0.641]	−0.005345 [−0.006326, −0.004403]	+0.02453 [0.02120, 0.02805]	+0.004494 [0.003411, 0.005586]
C3 − No Adapter	+4.852 [4.553, 5.142]	−0.02905 [−0.03352, −0.02464]	−0.03401 [−0.04214, −0.02591]	+0.04586 [0.04114, 0.05054]
C3 − GFH	−0.748 [−0.885, −0.609]	+0.01070 [0.008958, 0.012513]	−0.04553 [−0.04937, −0.04167]	−0.01881 [−0.02084, −0.01683]
C3 − Generic Linear	−0.593 [−0.837, −0.351]	+0.01351 [0.01160, 0.01549]	−0.02810 [−0.03197, −0.02439]	−0.01069 [−0.01300, −0.00850]
LLH − C3	+0.621 [0.391, 0.855]	−0.01214 [−0.01386, −0.01048]	+0.01658 [0.01311, 0.02003]	+0.009033 [0.006867, 0.011296]
LLH − Generic Linear	+0.028 [−0.015, 0.070]	+0.001372 [0.000762, 0.002081]	−0.01152 [−0.01337, −0.00996]	−0.001660 [−0.002096, −0.001256]

These results show sensitivity to the tested stage allocation, not a universally preferred placement or an equal-realized-dose effect.

Fresh B is a separate $N = 150$ cohort and reports LLH against LFM-EXACT, not C3. The paired LLH–LFM-EXACT differences are +0.325 dB FG-PSNR (95% CI [+0.209, +0.446]) and −0.00460 FG-LPIPS (95% CI [−0.00563, −0.00357]); both intervals fall within the recorded practical margins of ±0.5 dB and ±0.01. The saved lock states a pre-run freeze time, but the earliest preserved commit containing it is later than that time. We therefore report Fresh B as retrospective sensitivity rather than authenticated preregistered confirmation.

A separate Core-7 summary uses the same 276 object labels and recorded checkpoint lineage as Fresh C but does not establish common runner or complete input-tensor identity across its seven conditions. Its post-hoc C3–GFH contrasts point in opposite directions for Full-view PSNR (+0.725229 dB, 95% CI [+0.647134, +0.806467]) and FG-PSNR (−1.214235 dB, 95% CI [−1.313952, −1.114705]). It remains a retrospective diagnostic, not a same-runner confirmation or reproduction of the original +0.96 dB result (Supplementary Section S5).

4.6. Layer-by-Window Sensitivity

The A2 and A3b maps test response heterogeneity over three adapter groups and five denoising windows. Their object-cluster statistics describe variation within the tested interventions; they are not effect estimates for a dose-independent schedule mechanism.

A3b has no common-support rows across its three layer groups among 2,250 layer-step observations. The within-group central 90% realized-dose ranges are 56,359–79,927 (deep), 11,421–19,336 (middle), and 2,816–5,193 (shallow); they do not overlap. Dose-adjusted extrapolations cannot supply that missing support. Together with the Fresh B lock chronology, this limits A3b to retrospective sensitivity evidence. The A2/A3b maps support measurable response heterogeneity under their tested regimes, not equal-dose layer-by-time causality or an optimal schedule.

A separate retrospective analysis on the `FRESH_CONFIRM_300` cohort found nine of ten plus-one-corrected associations between LLH–GFL image-space deltas and ground-truth descriptors; the FG-LPIPS/coverage test had adjusted $p = 0.0508$. FG-PSNR means ranged from +4.140 dB in the lowest Laplacian-variance quartile to −0.087 and −0.029 dB in the upper two quartiles. These are cohort-specific associations, not a causal complexity effect, a threshold, a schedule selector, or a texture-fidelity result.

4.7. Appearance and Color Failure Analysis

The selected color diagnosis contains 26 failure-enriched objects: six earlier figure cases and 20 Fresh C panels. It is not a random sample, a new holdout, or a population estimate. No Adapter, GFL, and LLH were compared on the same inputs, checkpoint, and runner with six target views per object (78 object-condition records). The following paired object-bootstrap intervals are descriptive (20,000 resamples; seed 61008), and no confirmatory multiplicity family was defined.

GFL lowered reference-relative CIEDE2000 in 24/26 objects, while pink or purple shifts remained in the two cases shown in the submitted 01549 Figure 4 under No Adapter, GFL, and LLH. The decoded RGB output showed the shift before UV projection or baking, and a late denoising snapshot (update 40 of 50) showed it as well. In one selected case, a ground-truth VAE encode–decode isolate had CIEDE2000 4.086 and $\Delta a^* \approx -1.35$, compared with 22.032 and +1.977 for GFL. This weakens ordinary VAE round-trip error as the sole explanation in that case; it does not identify the complete cause or rule out other model and conditioning effects.

LLH–GFL also increased mean CIEDE2000 while reducing foreground gradient magnitude (−0.06772), high-frequency

Table 5. Post-inspection strict-276 placement analysis. Entries are C3 minus temporal L–L–H with pointwise 95% object-bootstrap intervals; positive PSNR/SSIM and negative LPIPS differences favor C3. Temporal L–L–H is favored on all six formal endpoints. The analysis is separate from Fresh C, and its scalar sums do not match realized residual dose.

Endpoint	Mean difference [95% CI]	Direction
Full-view PSNR (dB)	−0.268100 [−0.311101, −0.227797]	L–L–H
FG-PSNR (dB)	−0.177084 [−0.230539, −0.130784]	L–L–H
FG-SSIM	−0.008851 [−0.011455, −0.006291]	L–L–H
Edge-SSIM	−0.035301 [−0.036581, −0.034031]	L–L–H
FG-LPIPS	+0.001896 [+0.001205, +0.002578]	L–L–H
Full-view LPIPS	+0.008658 [+0.007465, +0.009831]	L–L–H

Table 6. Object-cluster tests of the observed layer-by-window response maps. A2 is exploratory; A3b is retrospective Fresh B sensitivity because the preserved history cannot authenticate pre-run availability of its lock. Interaction share and cell RMSE summarize the observed map, not a causal effect.

Map, cohort	Endpoint	Wald $W(8)$	Reported p	Share; cell RMSE
A2, $N = 300$	FG-PSNR	1879.963	$< 10^{-300}$	56.8%; 0.1222 dB
A2, $N = 300$	FG-LPIPS	481.081	8.04×10^{-99}	43.5%; 0.001695
A3b, $N = 150$	FG-PSNR	525.700	3.43×10^{-107}	36.0%; 0.0205 dB
A3b, $N = 150$	FG-LPIPS	225.015	5.37×10^{-43}	14.0%; 0.000342

Table 7. Selected paired results from the failure-enriched 26-object color diagnosis. Negative CIEDE2000 differences indicate lower reference-color error for the first condition. The set is descriptive and does not estimate population prevalence.

Contrast	Mean CIEDE2000 difference [95% CI]	Objects
GFL – No Adapter	−19.654 [−26.370, −13.248]	24/26 lower
LLH – GFL	+3.583 [+0.462, +6.712]	–

energy (−0.00887), and color entropy (−0.598). Lower local variation therefore did not imply better reference-color agreement or greater texture fidelity. The terminology follows Section 3.2: texture variation, image-space agreement, reference-color agreement, subjective preference, baked 3D appearance, and PBR/material fidelity are reported separately.

Three tested interventions failed to validate a color correction. The C1 chroma-anchor candidate passed its four-object numeric gate (mean unseen-view ΔFG-CIEDE2000 = −1.5602, 95% CI [−2.2621, −0.6299]) but failed locked visual review, including a cyan/blue rim. B2's four-object numeric result was +0.1729 (95% CI [−0.2044, +0.6810]), with only one of four objects improving and no clear visual correction. A feature-relative residual gate increased mean CIEDE2000 versus C3 by 1.35 across 26 objects and by 3.18 among the six earlier cases, with no clear visual reduction in the pink shift. These are sensitivity and failed-candidate results; no color-cast repair is validated.

The historical $N = 20$ GLB-native check used 20 UV-compatible objects from the original 24-object pool and 11 unseen views, rendered with unlit base color. C3–GFL FG-PSNR was −0.066 dB (95% CI [−0.413, +0.290]), FG-LPIPS was −0.004719 (95% CI [−0.007504, −0.002544]), and CIEDE2000 was +0.133 (95% CI [−1.161, +1.319]): endpoint directions are mixed, and these results do not establish full PBR/material fidelity or seam quality. The separate Fresh C 24-object bake did not pass its input gate: 9/24 assets had out-of-range native UVs, and 5 assets/27 primitives used unsupported BLEND alpha. No valid Fresh C bake or seam results were generated. A native-default adapter comparison is also absent; No Adapter at zero scale is not the native-default path.

Fresh C full-object and same-ROI examples.. The submitted 26-object texture audit and its quantitative Figure 5 cannot be re-traced to object IDs, selection criteria, raw values, and an executable metric definition in the current evidence archive. We therefore remove its 92% prevalence statement, its "true shape gain" classification, and its RGB-standard-deviation/gradient ratios from the scientific claims. Those statistics would describe variation, not texture fidelity, even if their source records were complete.

Figure 7 restores the original Fig. 4 role of showing full-object appearance beside a clearly marked local texture region. It compares the reference with GFL, GFH, and C3 for one selected C3–GFL favorable case and one adverse case; every column uses the same target view and pixel ROI within its object. Figure 4 retains the original C3-versus-fixed-scale question and adds a separate Generic Linear counterexample. These are selected Fresh C RGB illustrations, not a prevalence estimate or a texture-fidelity test. The full six-view, seven-condition panels, including No Adapter and LLH, remain in Supplementary Figure S1.

4.8. Cross-Interface Evaluation

Table 8 places the primary GeoTex path beside the auxiliary MV-Adapter and MVDiffusion interventions. The methods use different conditioning interfaces and scale semantics; object-level results are not pooled across systems.

For the 98-object MV-Adapter probe, the original additive-residual bootstrap had an invalid zero-power null and was not used for interaction inference. The exploratory object-cluster Wald test is reported with its corrected family; its non-detection does not establish absence of interaction. The 76-object mapping study is a layer mapping sensitivity analysis, not evidence that C3's temporal profile transfers. MVDiffusion's correspondence-aware path is also interface-specific. Taken together, these probes characterize tested interfaces and do not establish architecture-invariant scheduling.

4.9. First-round cluster-aware preference study

The 01549 first-round revision reported a blinded three-alternative forced-choice study on 30 objects with 24 participants. Participants compared conservative fixed-scale, aggressive fixed-scale, and C3 outputs for texture naturalness, shape

Table 8. Interface-specific intervention evidence. Nominal budgets are described within each implementation; no cross-interface realized-dose match is available. These rows are not architecture-matched replications.

System and injected signal	Control axis and budget	Cohort and principal result	Evidence grade and limitation
GeoTex-Adapter, Fresh C	Five-channel normal/depth/foreground adapter residuals added to UNet ResNet features; temporal requests with native layer caps. C3–GFL applied deep/middle sums differ (82.5 vs. 62.5); shallow sums are both 40.	N = 300. C3–GFL FG-PSNR is +0.522 dB [0.404, 0.641], but GFH and Generic Linear have more favorable means than C3 on all four displayed endpoints.	Matched, same-runner Fresh C image-space comparison; identifies the named paired contrast, not timing alone, a schedule optimum, or baked material quality.
MV-Adapter, layer/window probe	Condition guider plus decoupled multi-view attention; one of four down-block points is raised from 0.75 to 1.00 in one of five 10-step windows. Group residual-dose ratio is 1:0.57:1.46; no equal-realized-dose map was run.	98/99 planned objects. The exploratory four-metric Holm-corrected cluster-Wald family found no interaction detected after correction (primary FG-LPIPS Holm p = 0.303).	Interface-specific layer/window sensitivity; not C3 temporal transfer. A corrected non-detection does not show no interaction or equivalence.
MV-Adapter, mapping sensitivity	Four down-block points with a layer-wise LLH intervention under the alternate contiguous partition B'; per-point multipliers [0.52018, 0.52018, 1.47982, 1.47982] have weighted mean 1.	N = 76. Endpoint responses were mixed; the difference between B' and the original partition was not resolved on the reported metrics.	Budget-neutral nominal mapping sensitivity only; not a temporal transfer test, dose-independent mapping law, or general backbone result. The companion M3 profile also changes nominal budget and is retrospective.
MVDiffusion, compatibility interface	Correspondence-aware decoder-side CP-Blocks; the tested multiplier is a product of a layer factor and temporal scale over five blocks. Official SD2-depth weights were unavailable, so all panel conditions used the same compatibility base.	N = 75, 12 target views, 50 steps. L–LLH vs. G–FL had lower CIEDE2000 (−1.373 [−2.375, −0.418]) but lower PSNR (−0.454 [−0.752, −0.160]) and FG-SSIM (−0.0194 [−0.0294, −0.0099]); LPIPS did not separate.	Within-interface, mixed endpoint comparison on a different conditioning path; no shared scale semantics, cross-system dose match, or architecture-causal inference.

Table 9. First-round 01549 3AFC results retained as reported. Percentages use 720 decisions per criterion; preference is not a measure of reference texture fidelity.

Method	Texture naturalness	Shape consistency	Overall quality
s = 1.25	332/720 (46.1%)	96/720 (13.3%)	176/720 (24.4%)
s = 2.50	104/720 (14.4%)	331/720 (46.0%)	126/720 (17.5%)
C3 (TCAS)	284/720 (39.4%)	293/720 (40.7%)	418/720 (58.1%)

consistency, and overall quality; condition order was randomized per trial. Each criterion had 720 first-choice decisions. Table 9 reports the aggregate counts; participant-level responses are not redistributed with this revision.

The reported analysis used a two-level participant/object bootstrap with 10,000 resamples. C3's overall-quality share was 58.1% (95% CI [50.1%, 65.8%]); its differences from the conservative and aggressive baselines were +33.6 points (CI [19.3, 47.2]) and +40.6 points (CI [29.4, 51.8]). The single-criterion advantages for texture naturalness and shape consistency had intervals crossing zero. This bounded preference result does not validate texture fidelity, unseen-view quality, or the Fresh C cohort.

4.10. FAC controlled re-examination reported in 01549

A controlled re-examination of Full Adaptive Correction (FAC) was previously reported, with the base model, VAE, and GeoTex-Adapter frozen and roughly 6×10^3 controller parameters trainable. The stated protocol used the enhanced latent denoising objective (foreground weight 1.7, edge weight 1.3, and a latent-SSIM term), AdamW at peak learning rate 10^{-4} with 100 warm-up steps, and 2,000 training steps on a training pool reported as 1,706 objects and disjoint from evaluation. It reported a 300-object paired evaluation: every trained configuration reduced FG-PSNR versus TCAS by 1.10–2.94 dB, with at most 48/300 wins; the untrained C3-envelope control on 12 objects had ΔFG-PSNR = −0.18 dB (95% CI [−0.42, +0.05]). The seven reported configurations and endpoint values are retained in Supplementary Table S11.

These FAC outcomes are retained as previously reported results. They are not presented as a new 1008 training run. The current archive audit found that a recovered configuration points to a 1,200-object list rather than the 1,706-object pool stated in 01549, and does not bind the training seed and run outputs in one manifest. We therefore disclose the remaining reproducibility gap and make no code-release claim; TCAS remains the sole training-free method.

5. Limitations

Fresh C is a held-out 300-object comparison whose conditions share one frozen checkpoint, runner, seed policy, inputs, metric implementation, and recorded native-cap semantics. The checkpoint is a controlled retraining and differs from the v1 path named for 01549 Table 4; the 01549 Table 9/FAC work uses a separate v2 strong-residual adapter. These historical v1/v2 differences are reported separately from the within-Fresh-C paired comparison. Fresh C does not numerically reproduce the old +0.96 dB Table 4 estimate. The historical maintable endpoint implementation is not recoverable for a numerical equivalence claim. Core-7 and the post-inspection strict-276 profile-placement analysis use the clean-v2 276-object list and the checkpoint SHA recorded for Fresh C, but are separate runs from Fresh C; Core-7 condition records do not establish one shared runner or complete input-tensor identity. They remain retrospective supporting diagnostics rather than confirmatory comparisons.

Texture variation and texture fidelity are distinct. The submitted RGB standard deviation, Laplacian variance, and gradient summaries are not re-used as fidelity measures; their source records and definitions are incomplete, and such metrics can rise with artifacts. Fresh C reports PSNR, LPIPS, and SSIM

endpoints but has no validated texture-fidelity measure. The separate retrospective FRESH_CONFIRM_300 analysis reports a 9/10 adjusted association family and quartile heterogeneity for LLH–GFL image-space deltas; it does not measure fidelity, establish causality, or predict which schedule should be used. The separate 26-object color diagnosis is a failure-enriched, paired analysis rather than a population confirmation: GFL reduced reference-relative CIEDE2000 versus No Adapter, but visible hue mismatch persisted in the two cases shown in the submitted 01549 Figure 4, and LLH reduced local texture measures. These results do not validate a color repair or treat texture variation as fidelity (Supplementary Section S13). The 01549 cluster-aware preference study is a separate subjective choice result, not a fidelity or 3D-quality endpoint. Fresh C has no validated baked-mesh, seam, or PBR measure. Historical N=20 unlit base-color renders provide only a small, mixed unseen-view result.

An exploratory four-object color-pathway audit found that changing the global appearance embedding can shift chroma. A separate Fresh B input-only audit found that stored embeddings matched neither of the two compared source-conditioning reconstructions within 0.03 for 150/150 inputs. This makes cache/source alignment a plausible contributor, not an established general cause of color drift. The C1 chroma-anchor candidate met its numeric gate (mean unseen-view ΔFG-CIEDE2000 = -1.5602, 95% CI $[-2.2621, -0.6299]$, four of four) but failed the locked visual gate: one case showed no clear gain, and another retained mismatch with a faint cyan/blue rim. B2's selected raw-source embedding substitution did not meet its numeric gate (mean object-level change in the lower-is-better ΔFG-CIEDE2000 endpoint = +0.1729, 95% CI $[-0.2044, +0.6810]$, one of four improved) and showed no clear visual correction. No independent Fresh B repair validation was performed. These development findings identify a control pathway and a plausible contributor, not a validated correction or a complete root cause.

A separate, failure-enriched 26-object color diagnostic compared No Adapter, GFL, and LLH outputs with paired six-view references. GFL minus No Adapter reduced mean foreground CIEDE2000 by 19.65 (95% paired object-bootstrap CI $[-26.37, -13.25]$; 24/26 objects), but the two cases shown in the submitted 01549 Figure 4 still showed visible pink or purple hue shifts under No Adapter, GFL, and LLH. LLH minus GFL increased mean CIEDE2000 by 3.58 (95% CI $[+0.46, +6.71]$) while reducing measured local gradient and high-frequency variation. The selected diagnostic set is not an independent confirmation cohort, and variation measures are not fidelity measures. Pink coloration was visible in a decoded RGB snapshot before baking; a ground-truth VAE round trip did not reproduce the same direction or magnitude of the selected GFL mismatch. These findings narrow the location and exclude simple explanations, but do not identify a complete color-generation cause or validate a repair (Supplementary Section S13).

The requested TCAS profile does not uniquely determine the applied scale or correction magnitude. Fresh C's native wrapper caps shallow requests at 0.8, and residual norms depend on the adapter output. The strict-276 schedule sums are only approximately matched and do not equalize realized dose; A3b has no cross-layer common support. These data do not identify a dose-independent mechanism, an optimal selector, or a schedule that dominates all endpoints and objects.

We withdraw the CAI optimality proposition and the earlier positive FAC extension claim; the negative FAC controlled re-examination and the cluster-aware 01549 preference study are retained as previously reported first-round results with their stated scope and provenance limits. The old pooled +0.96 dB headline and broad texture-retention claim are not used as current evidence. CLIP-IQA is not used to establish texture fidelity. The revised claim boundaries are documented in the changelog.

6. Conclusion

This study treats geometry-adapter strength as an observable inference control in a multi-view texturing pipeline. TCAS specifies a requested temporal profile on the existing GeoTex-Adapter path; C3, (1.25, 2.50, 1.25), is the recorded profile evaluated here, not a general scheduling rule.

On the matched 300-object Fresh C cohort, C3 exceeds GFL by +0.522352 dB on the registered FG-PSNR endpoint (95% CI [+0.404160, +0.640669]), while GFH and Generic Linear have more favorable means than C3 on all four displayed endpoints. The applied deep/middle scale sums differ (82.5 for C3 and 62.5 for GFL; both have 40 in shallow), so this contrast does not isolate timing from applied scale. A separate post-inspection strict-276 placement analysis favors L–L–H over C3 across six endpoints, but does not match post-scale residual norms or provide an independent replication.

The empirical contribution is a reproducible, task-specific account of how an existing geometry-residual control behaves under its native wrapper, evaluated against fixed and scheduled alternatives. The retrospective texture-complexity analysis and bounded layer/window maps show that responses vary across objects and interventions, but do not identify a causal complexity mechanism or an equal-dose scheduling law. Current evidence remains image-space; it does not establish texture fidelity, baked-mesh quality, seams, PBR appearance, a universal schedule, or general cross-backbone transfer.

Data availability

Aggregate results and the LaTeX source required to rebuild the manuscript and Supplement are included in the author-review package. Raw generated-image archives, trained checkpoints, source 3D assets whose redistribution rights are not established, and participant-level responses are not included. Additional analysis records may be requested from the corresponding author, subject to licensing, privacy, and archive constraints.

Declaration of competing interests

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

References

[1] Liu R, Wu R, Van Hoorick B, et al. Zero-1-to-3: Zero-shot one image to 3D object. *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2023.

[2] Shi R, Chen H, Zhang Z, et al. Zero123++: A single image to consistent multi-view diffusion base model. arXiv:2310.15110, 2023.

[3] Liu Y, Lin C, Zeng Z, et al. SyncDreamer: Generating multiview-consistent images from a single-view image. *International Conference on Learning Representations*, 2024.

[4] Long X, Guo Y C, Lin C, et al. Wonder3D: Single image to 3D using cross-domain diffusion. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024.

[5] Shi Y, Wang P, Ye J, et al. MVDream: Multi-view diffusion for 3D generation. *International Conference on Learning Representations*, 2024.

[6] Wang P, et al. ImageDream: Image-prompt multi-view diffusion for 3D generation. arXiv:2312.02201, 2023.

[7] Xu J, et al. InstantMesh: Efficient 3D mesh generation from a single image with sparse-view large reconstruction models. *Advances in Neural Information Processing Systems*, 2024.

[8] Li P, Liu Y, Long X, et al. Era3D: High-resolution multiview diffusion using efficient row-wise attention. *Advances in Neural Information Processing Systems*, 2024.

[9] Chen D Z, et al. Text2Tex: Text-driven texture synthesis via diffusion models. *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2023.

[10] Richardson E, Metzer G, Alaluf Y, et al. TEXTure: Text-guided texturing of 3D shapes. *ACM SIGGRAPH 2023 Conference Proceedings*, 2023.

[11] Poole B, Jain A, Barron J T, Mildenhall B. DreamFusion: Text-to-3D using 2D diffusion. *International Conference on Learning Representations*, 2023.

[12] Chen R, Chen Y, Jiao N, Jia K. Fantasia3D: Disentangling geometry and appearance for high-quality text-to-3D content creation. *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2023.

[13] Shao M, Xiong F, Sun Z, Xu M. MVPainter: Accurate and detailed 3D texture generation via multi-view diffusion with geometric control. arXiv:2505.12635, 2025.

[14] Kynkäänniemi T, Aittala M, Karras T, et al. Applying guidance in a limited interval improves sample and distribution quality in diffusion models. *Advances in Neural Information Processing Systems*, 2024.

[15] Perla S R K, Zhang H, Mahdavi-Amiri A. Advances in neural 3D mesh texturing: A survey. *Computer Graphics Forum*, 2026, e70392.

[16] Guzmán J E, Mould D, Paquette E. Adaptive multiresolution exemplar-based texture synthesis on animated fluids. *Computers & Graphics*, 2026, 134: 104490.

[17] Lv X, Wu Z, Xu J, et al. Interpretable procedural material graph generation via diffusion models from reference images. *The Visual Computer*, 2025, 41(13): 11195–11205.

[18] Zeng X, Chen X, Qi Z, et al. Paint3D: Paint anything 3D with lighting-immersive texture generation. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024.

[19] Cheng W, Mu J, Zeng X, et al. MVPaint: Synchronized multi-view diffusion for painting anything 3D. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2025.

[20] Zhao Z, Lai Z, Lin Q, et al. Hunyuan3D 2.0: Scaling diffusion models for high resolution textured 3D assets generation. arXiv:2501.12202, 2025.

[21] Huang Z, et al. MV-Adapter: Multi-view consistent image generation made easy. *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2025.

[22] Tang S, Zhang F, Chen J, Wang P, Furukawa Y. MVDiffusion: Enabling holistic multi-view image generation with correspondence-aware diffusion. *Advances in Neural Information Processing Systems*, 2023.

[23] Liu Y X, Xie M S, Liu H Y, Wong T T. Text-guided texturing by synchronized multi-view diffusion. *SIGGRAPH Asia 2024 Conference Papers*, 2024.

[24] Xu Q, Han X G, Zhang L. WonderTex: Consistent-and-seamless texture generation with text-guided multi-view image diffusion models. *IEEE Transactions on Visualization and Computer Graphics*, 2026, 32(7): 5815–5825.

[25] Mou C, Wang X, Xie L, et al. T2I-Adapter: Learning adapters to dig out more controllable ability for text-to-image diffusion models. *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024.

[26] Zhang L, Rao A, Agrawala M. Adding conditional control to text-to-image diffusion models. *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2023.

[27] Ye H, Zhang J, Liu S, Han X, Yang W. IP-Adapter: Text compatible image prompt adapter for text-to-image diffusion models. arXiv:2308.06721, 2023.

[28] Hu E J, Shen Y, Wallis P, et al. LoRA: Low-rank adaptation of large language models. *International Conference on Learning Representations*, 2022.

[29] Liu S Y, Wang C Y, Yin H, et al. DoRA: Weight-decomposed low-rank adaptation. *International Conference on Machine Learning*, 2024.

[30] Zhang Q, Zuo S, Liang C, et al. AdaLoRA: Adaptive budget allocation for parameter-efficient fine-tuning. *International Conference on Learning Representations*, 2023.

[31] Kulkarni A S, et al. Scheduled Style Injection: Expanding the Style-Content Pareto Frontier in Training-Free Diffusion-Based Style Transfer. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), NTIRE Workshop*, 2026.

[32] Ho J, Jain A, Abbeel P. Denoising diffusion probabilistic models. *Advances in Neural Information Processing Systems*, 2020.

[33] Song J, Meng C, Ermon S. Denoising diffusion implicit models. *International Conference on Learning Representations*, 2021.

[34] Ho J, Salimans T. Classifier-free diffusion guidance. *NeurIPS Workshop on Deep Generative Models and Downstream Applications*, 2021. arXiv:2207.12598.

[35] Castillo A, Kohler J, Pérez J C, et al. Adaptive guidance: Training-free acceleration of conditional diffusion models. *Proceedings of the AAAI Conference on Artificial Intelligence*, 2025.

[36] Wang Z, Bovik A C, Sheikh H R, Simoncelli E P. Image quality assessment: From error visibility to structural similarity. *IEEE Transactions on Image Processing*, 2004, 13(4): 600–612.

[37] Zhang R, Isola P, Efros A A, Shechtman E, Wang O. The unreasonable effectiveness of deep features as a perceptual metric. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2018.

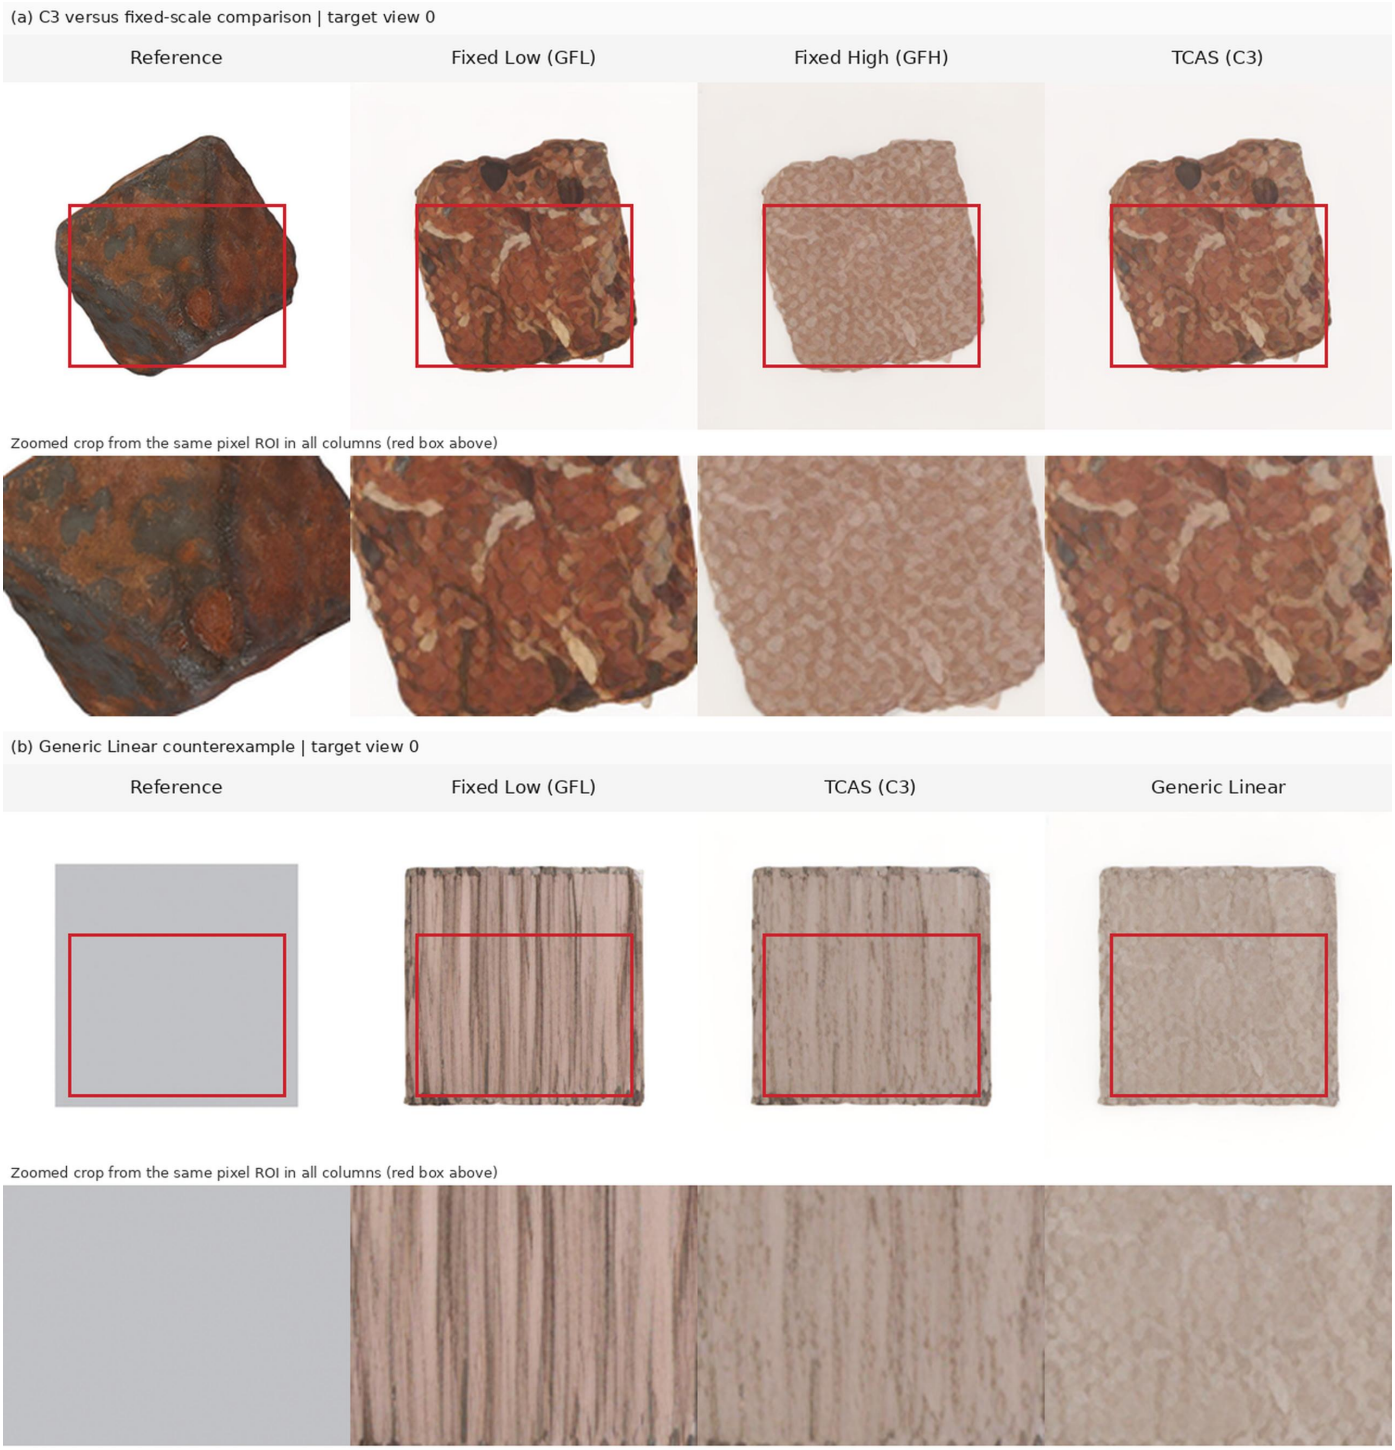


Fig. 4. Matched Fresh C comparisons under the same frozen object, reference, camera view, checkpoint, runner, seed policy, and inference protocol. (a) Reference, fixed-low GFL, fixed-high GFH, and C3 preserve the original fixed-scale comparison question; the selected object illustrates an endpoint trade-off between C3 and GFH. (b) Reference, GFL, C3, and Generic Linear show a counterexample in which both object-level FG-PSNR (C3 minus Generic Linear = −3.491 dB) and FG-LPIPS (= +0.08172) favor Generic Linear. Each context image marks the same pixel ROI used for the crop below. The examples were selected by outcome, so they do not estimate prevalence or establish a general winner or texture fidelity. They show qualitative local appearance rather than colorimetric agreement; object-specific hue mismatch may remain. Full six-view comparisons remain in Supplementary Figure S1. Attribution is in Supplementary Table S1; input and ROI hashes are in the Fresh C figure manifest.

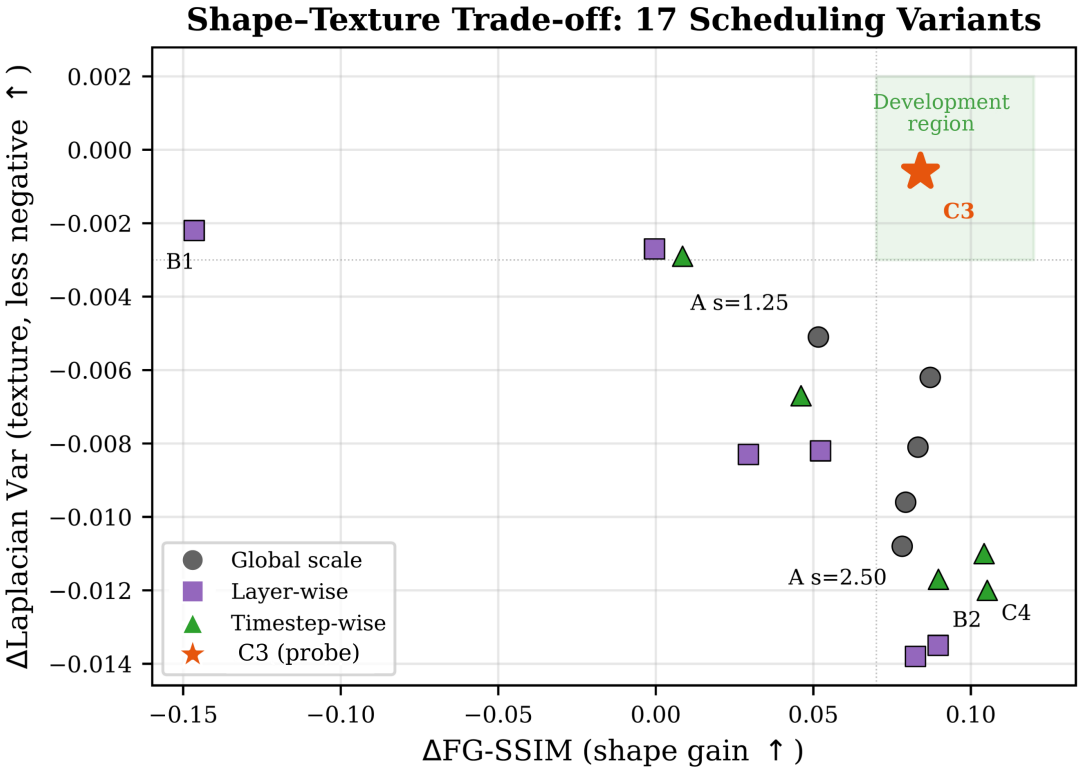


Fig. 5. Historical 17-variant shape–texture development plot from the submitted 01549 paper, restricted to the 24-object development probe and not an independent validation set. The plot labels have been locally changed to "Development region" and "C3 (probe)" to identify their development-set role; neither the shaded region nor C3 denotes a unique optimum, a texture-fidelity ranking, or a general schedule law. The underlying row-level candidate outputs are not available for independent reanalysis.

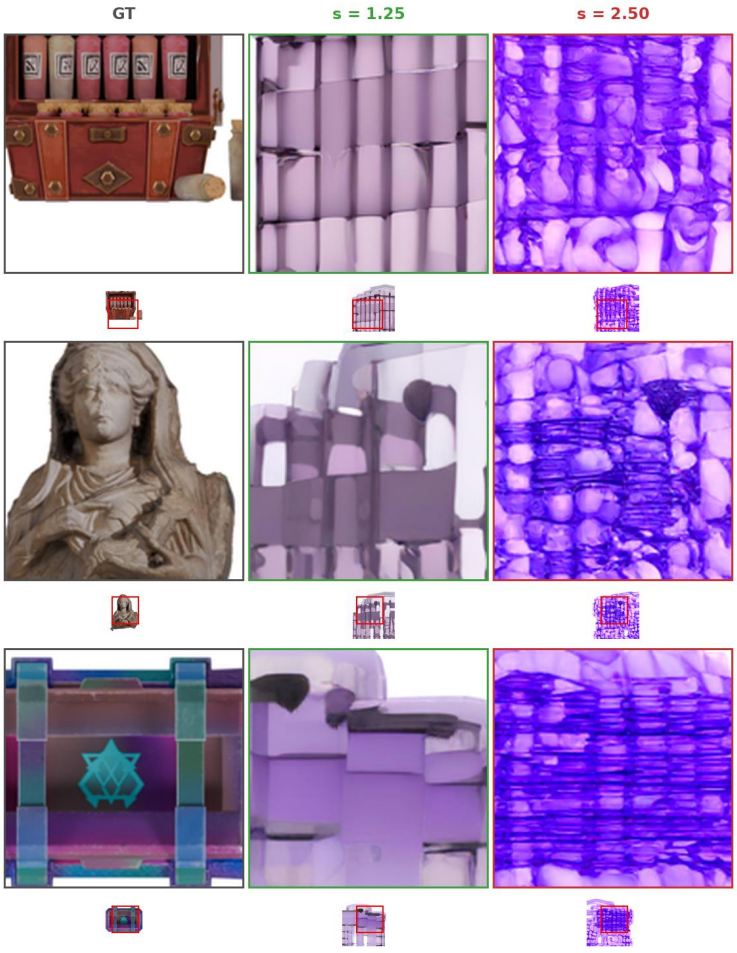


Fig. 6. Historical color-shift examples reproduced from the authenticated 01549 Figure 4. Fixed-scale target-view outputs show visible pink/purple mismatch against the reference. The original panel-level object/run mapping is not recoverable. This figure documents the historical failure that motivated the diagnosis; it is not a Fresh C panel, a representative sample, or a colorimetric estimate.

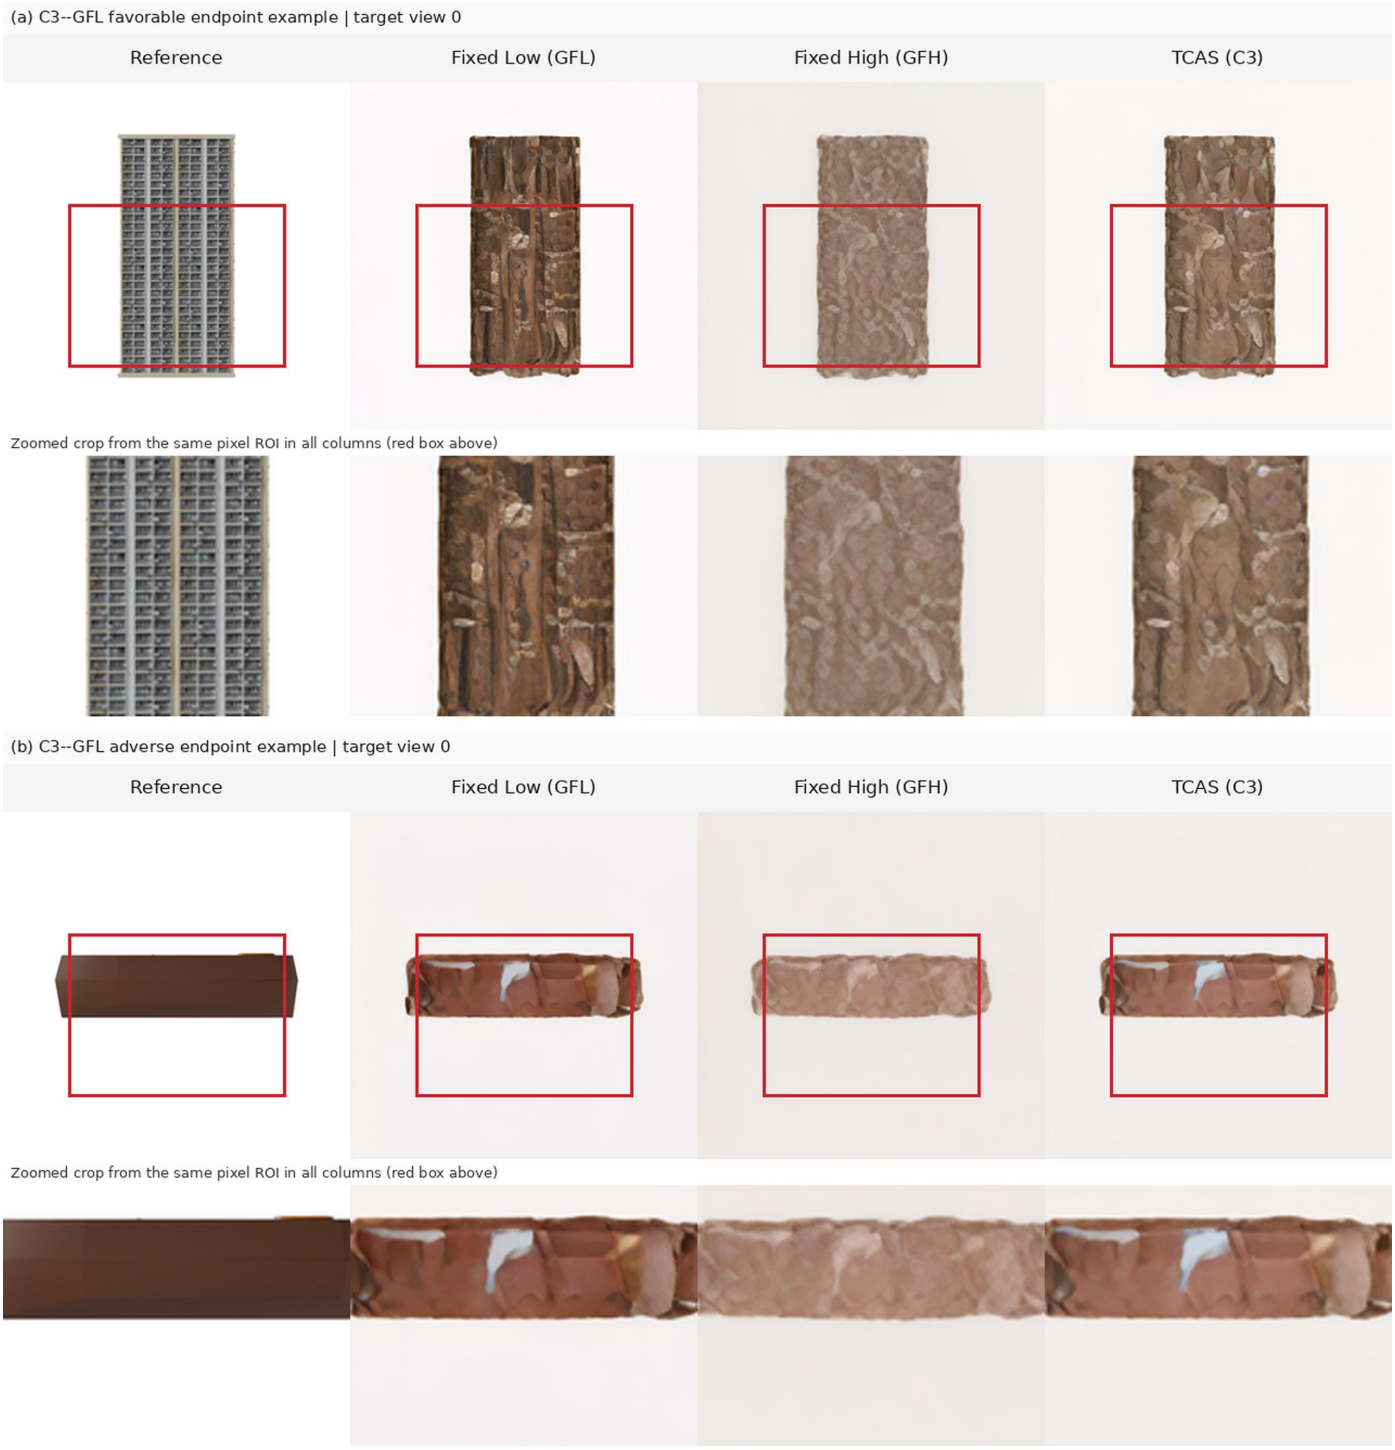


Fig. 7. Matched Fresh C comparisons under the same frozen object, reference, camera view, checkpoint, runner, seed policy, and inference protocol. Columns show the reference, fixed-low GFL, fixed-high GFH, and TCAS C3. Each panel shows the full target-view context with the shared ROI marked in red, followed by the same-pixel crop. Case (a) is selected because the paired C3–GFL FG-PSNR and FG-LPIPS differences both favor C3; case (b) is an adverse case in which both endpoints favor GFL. The outcome-selected examples are illustrative only and do not estimate case frequency or texture fidelity. They show qualitative local appearance rather than colorimetric agreement; object-specific hue mismatch may remain. The shown images are target-view RGB outputs, not baked textures, unseen-view renders, PBR results, or human-preference evidence. Model credits are in Supplementary Table S1; source-image hashes and ROI coordinates are in the Fresh C figure manifest.