Research highlights — revision draft

- Equal-budget tests on 300 objects reject a universally preferred middle window.
- Native MV-Adapter tests on 50 objects reveal endpoint-dependent responses.
- A calibrated ten-object bake pilot supports a small late-high LPIPS benefit.
- C3 counter-results and persistent appearance failures limit deployment claims.
- Object-paired records, calibration tests and analysis scripts are supplied.