Reviewer risks and experiment decisions

10 October 2026 · Author review only · `SUBMISSION_APPROVED=NO` · No new GPU run launched

Review-source note: the exact local plain-text source is retained at `1008/audits/source_evidence/r2/reviewer_materials/SECOND_ROUND_REVIEWER_COMMENTS.txt` (SHA-256 `52d5a02718cbcf1b452293fe5a6953231f10fc9ce4cdf6cce400e86deb521245`). Its text matches the supplied attachment after line-ending normalization. The separately preserved PDF text extraction (SHA-256 `1d2efb8f9250f8162304be57ce1d71c6d654d35511a186ad6e3c9b869ace5af5`) contains pagination and extraction artifacts. The source itself says "Initial submission"; the author-confirmed round map records this packet as the second-round review of 01549.

Author checks: verify final page references and Figure 6 attribution.

Issue	Evidence and remaining risk	Decision
R1.1 Color / 3D fidelity — high, partial	Failure-enriched N=26 diagnosis; no validated color repair. No Adapter is not native-default. Fresh C bake gate failed (9/24 out-of-range UV inputs; 5 assets / 27 primitives use unsupported BLEND). Historical N=20 unlit check is mixed. Current Figure 6 has confirmed CC BY 4.0 rights and a hash-linked single-case selection record; it is not a prevalence estimate or one of the unmapped legacy cases.	**E1: highest priority, no-go now.** Do a CPU-only preflight on an unchanged, frozen in-cohort asset (estimated 1–2 h). Require valid original UVs/materials, sampler/import-export integrity, finite output, and repeatable pixels. If it passes, a small native-default / No Adapter / GFL / C3 bake and unseen-view/seam pilot is estimated at 0.5–1 day. Failure means stop, no substitution or input repair.
R1.2 / R1.4 Matched result and uncertainty — bounded	Fresh C: N=300; C3−GFL FG-PSNR +0.522352 dB (95% CI [+0.404160,+0.640669]). GFH and Generic Linear have stronger means on all four shown endpoints. Applied deep/middle sums differ, 82.5 vs. 62.5.	No new experiment. Keep the positive endpoint result and counter-results together; no timing-only, equivalence, or non-inferiority claim.
R1.3 Texture fidelity — partial	Variation proxies and subjective preference are not fidelity. Fresh C has no validated reference-texture-fidelity endpoint.	No experiment selected. A future study needs a defined reference, held-out protocol, and validated endpoint before data collection.
R1.5 FAC — partial	Reported training pool is 1,706 objects; recovered config points to 1,200. Seed, output, and runtime are not bound in one manifest.	Retain as prior negative evidence with the mismatch disclosed; make no full-reproduction or public-code claim.
R2.1 Contribution — editorial risk	SSI establishes prior layer/time scheduling, including geometry-conditioned ControlNet strength. TCAS is a task-specific GeoTex-Adapter residual intervention with measured cap, dose, and endpoint limits.	No schedule sweep. Present an empirical graphics-system contribution; contribution sufficiency remains an editor judgment.

Issue	Evidence and remaining risk	Decision
R2.3 Cross-backbone — high, open	The 76-object MV-Adapter layer/time panel has 228 new rows, but its feature-scaling path differs from GeoTex residual correction. Its protocol and results first appear in the same preserved commit, all 30 intervals are pointwise/unadjusted, and wrapper identity covers only 9/9 cases across three objects, not the official native scale-1 path. MVDiffusion's 75-object/525-row panel used a compatibility base and scored six target views within each 12-view output grid. The corrected raw L−LLH−G−FL CIEDE2000 is +1.373 [+.418,+2.375] (higher error); the archive's negative value was improvement-coded. In P5, raw G−LLH−G−LHL FG-LPIPS is −0.001371 [−0.002642,−0.000141], favoring G−LLH; the archived report's verbal direction was wrong. A separate mirror check matched 1,552 tensors and 12 views for one object, but did not rerun the panel.	**E2: no-go.** These interface probes do not meet the matched-transfer gate. A future pilot would first need a verified common scale meaning and within-backbone native/fixed/scheduled controls; no GPU run is authorized or queued. Keep current results interface-specific and do not infer equivalence or transfer.
Dose / temporal placement — open	Strict-276 L–L–H beats C3 on six endpoints, but its layer-wise post-scale corrections differ; A3b's raw-L_2 cross-layer support test is dimension-confounded, so valid dose support is unresolved. A 50-step CPU profile preflight passed (82.5 deep/middle and 40 shallow applied-scale sums), but this is not realized dose. No inference: GPU-idle gate failed (~24/32 GiB on both RTX 5090s). The base-model directory was not content-hashed.	**E3: no-go.** Reconsider only with a frozen design using a dimension-aware dose definition and a useful estimand. Any GPU pilot follows only if the design is locked and resources are cleared. The current evidence does not identify an equal-dose or dose-independent causal effect.