


TCAS: stage-wise low-high-low schedule



$C3 = (1.25, 2.50, 1.25)$: training-free, from measured stage utilities

Evaluation: 300 objects + disjoint holdout

Uniform high scale: structure up, texture down
(Lap. var. 67.8%, RGB std 64.3% of conservative)

C3: +0.96 dB PSNR vs. $s=2.50$, FG-SSIM comparable (n.s.); texture retained (Lap. var. **83.5%**)

58.1% blinded first choices (cluster bootstrap CI excludes zero)

No further search: 24-object probe $\rightarrow$ disjoint
276-object holdout

Timestep-Conditioned Adapter Scaling (TCAS) for Multi-view Diffusion Texture Generation