



A — Uniform adapter scale: structure up, texture down B — TCAS: stage-wise low-high-low schedule



high scale $\Rightarrow$ flat surfaces, detail loss

C3 = (1.25, 2.50, 1.25): training-free, derived from measured stage utilities (CAI)

C — Evaluation: 300 objects + disjoint holdout

No further search: 24-object probe $\rightarrow$ disjoint 276-object holdout

+0.96 dB PSNR vs. aggressive uniform scaling $s=2.50$

FG-SSIM comparable to $s=2.50$ (n.s.)

Texture retention vs. $s=1.25$:

Lap. var. **83.5%**, RGB std **97.0%**, grad. **88.3%**
($s=2.50$: 67.8%, 64.3%, 80.7%)

58.1% blinded overall first choices
(cluster bootstrap CI excludes zero)

TCAS: plug-and-play, no retraining, negligible overhead

Timestep-Conditioned Adapter Scaling (TCAS) for Multi-view Diffusion Texture Generation